

Theoretical Proposal for Discussion

Formalizing Cross-Scale Interactions in Autonomous Agents: Perspective Analysis of the Global State Vector

Authors

Timur Urmanov¹, Kamil Gadeev², Bakhtier Iusupov³

¹ Conceptualization, Formal Analysis, Methodology, Writing

² Software, Methodology, Validation, Ontological Framework

³ Supervision, Project Administration, Integration

Correspondence

¹ urmanov.t@gmail.com

² gadeev.kamil@gmail.com

³ usupovbahtiayr@gmail.com

Target Audience

Verses.ai Research Team

Active Inference Community

Multi-Agent Systems Researchers

October 2025

License: CC BY-NC 4.0 (Attribution — Non-commercial)

Abstract

Current multi-scale control architectures for autonomous AI agents, such as the Global State Vector 2.0 (GSV 2.0), rely on heuristic cross-coupling terms to manage interactions between strategic control variables. While functionally effective, these terms lack first-principles derivation, limiting theoretical integrity and interpretability.

We resolve this critical gap by demonstrating that these cross-coupling terms are not arbitrary heuristics but represent the mathematically necessary rotational dynamics for maintaining adaptive, non-equilibrium behavior. Through three complementary theoretical formalisms, we prove this claim:

First, employing the Helmholtz-Hodge decomposition, we rigorously show that the cross-terms constitute the system’s *rotational flow component*, characterized by non-zero curl ($\nabla \times \mathbf{r} \neq 0$). This rotational flow is mathematically essential for breaking detailed balance and maintaining a non-equilibrium steady state (NESS)—the hallmark of adaptive, life-like systems.

Second, through information geometry, we demonstrate that these couplings emerge naturally from off-diagonal structure in the Fisher Information Matrix (FIM), reframing GSV dynamics as natural gradient descent where couplings reflect inherent statistical dependencies between meta-parameters.

Third, we provide a mechanistic interpretation, showing how cross-terms function as multiplicative gates analogous to neuromodulatory mechanisms in biological neural circuits, dynamically regulating information flow between strategic modules.

By unifying physical necessity (NESS maintenance), statistical optimality (natural gradients), and computational mechanism (gating), this work elevates GSV’s cross-terms from functionally motivated additions to principled, mathematically grounded components of a robust AI control architecture. Detailed numerical validation confirms theoretical predictions with quantitative metrics and confidence intervals.

Contents

1	Introduction: The Problem of Heuristic Coupling in Multi-Scale Agent Architectures	4
1.1	The Challenge of Multi-Scale Adaptation	4
1.2	The Global State Vector (GSV 2.0) as a Proposed Solution	4
1.3	Identifying the Theoretical Gap	4
1.4	Thesis and Roadmap	5
2	The Global State Vector: A Stochastic Dynamical System for Strategic Control	6
2.1	Mathematical Specification	6
2.2	Analysis of System Components	6
2.2.1	Drift Terms: Drivers and Decay	6
2.2.2	Stochastic Terms	7
2.2.3	Cross-Coupling Terms	7
3	Formalism I: Helmholtz Decomposition and the Role of Rotational Flows	8
3.1	The Helmholtz-Hodge Decomposition	8
3.2	Decomposing the GSV Drift Vector	9
3.2.1	The Gradient Flow Component	9
3.2.2	The Rotational Flow Component	10
3.3	Functional Interpretation: NESS and Breaking Detailed Balance	11
4	Formalism II: Information Geometry and Natural Gradient Dynamics	13
4.1	Information Geometry and the Fisher Information Metric	13
4.2	Natural Gradient Descent	14
4.3	The Role of a Non-Diagonal FIM	15
4.4	The GSV Dynamics as Natural Gradient Flow	15
4.5	Practical Considerations: Regularization and Robustness	17
5	Formalism III: Strategic Gating as a Neuromodulatory Analogue	17
5.1	Multiplicative Gating in Recurrent Neural Networks	18
5.2	The tanh Function as a Probabilistic Gate	18
5.3	GSV Cross-Terms as Strategic Gates	18
5.4	Analogy to Biological Neuromodulation	19
6	Synthesis and Discussion: Unifying Perspectives on Systemic Coupling	20
6.1	Three Lenses, One Truth	20
6.2	The Unified Picture	21
6.3	Implications for AI Design and Convergent Evolution	22
7	Conclusion and Future Directions	23
7.1	Recapitulation: From Heuristic to Principled	23
7.2	Implications: Toward a Universal Theory of Multi-Scale Control	23
7.2.1	Prescriptive Power	23
7.2.2	Convergent Evolution and Universal Principles	24
7.3	Future Research Directions	24

7.3.1	Automated Discovery of Coupling Structure	24
7.3.2	Extension to Higher-Dimensional Systems	24
7.3.3	Connection to Active Inference and the Free Energy Principle	25
7.3.4	Biological Validation	25
7.4	Final Reflection	25
8	Glossary	26
9	References	27

1 Introduction: The Problem of Heuristic Coupling in Multi-Scale Agent Architectures

1.1 The Challenge of Multi-Scale Adaptation

A fundamental challenge in the development of truly autonomous artificial intelligence is enabling agents to operate and adapt effectively over long temporal horizons. While significant progress has been made in reactive cognition — decision-making on the order of milliseconds to seconds — genuine autonomy requires strategic adaptation at much slower timescales, from minutes to days. Consider an agent deployed in a dynamic environment: a strategy of high exploration may be optimal in a novel, safe context but becomes catastrophically maladaptive when the environment shifts to a volatile, resource-scarce regime. The core problem is not what action to take next, but rather what *mode of operation* to inhabit. How should an agent autonomously modulate its own high-level learning and decision-making parameters — its meta-strategy — based on accumulated experience?

This requirement for hierarchical, multi-scale control is not an idiosyncrasy of artificial systems but appears to be a universal principle of complex adaptive systems. Biological systems provide the canonical example. The brain operates across numerous, nested timescales, from the rapid firing of neurons to the slower dynamics of working memory and, at the longest scales, the neuromodulatory control exerted by hormones like cortisol and dopamine. This slow layer does not dictate specific actions but rather modulates the parameters of the fast cognitive machinery, adjusting sensory gains, decision thresholds, and behavioral proclivities like risk aversion or novelty-seeking. This architectural motif is thought to be a product of convergent evolution, an obligatory design for any system that must balance rapid responses with long-term strategic integrity in a complex world.

1.2 The Global State Vector (GSV 2.0) as a Proposed Solution

To address this challenge, our previous work (Urmanov, Gadeev, Iusupov, 2025, DOI: [10.5281/zenodo.17286613](https://doi.org/10.5281/zenodo.17286613)) introduced the **Global State Vector 2.0 (GSV 2.0)**, a formal framework for multi-scale control in artificial agents. It posits a low-dimensional, continuous dynamical system that evolves on a slow timescale, modulating the parameters of a fast-timescale cognitive architecture. The GSV is defined by four state variables, or axes, each representing a fundamental trade-off that any autonomous agent must manage:

1. **Arousal** (S_A): The mobilization of resources in response to threat or uncertainty.
2. **Exploration/Exploitation** (S_E): The balance between seeking novel information and leveraging known strategies.
3. **Plasticity** (S_P): The rate of change in the agent’s own architecture or fundamental parameters.
4. **Social Adaptation** (S_S): The degree of coordination and alignment with other agents.

By modeling these strategic dimensions as a set of coupled Stochastic Differential Equations (SDEs), the GSV 2.0 framework provides an explicit, interpretable control layer that adapts an agent’s meta-strategy based on continuous feedback from its performance and environment.

1.3 Identifying the Theoretical Gap

While our previous work established a strong connection between the primary dynamics of the system (the driving and decay terms for each axis) and the Free Energy Principle, we also transparently

identified a crucial limitation: the cross-coupling terms that model the influence of one axis upon another. As stated in that paper, "*we acknowledge that the cross-coupling terms are motivated by functional necessity rather than strict derivation from first principles.*" This admission highlights a central theoretical gap: while terms like the one coupling Arousal to Exploration are functionally intuitive and essential for producing adaptive behavior, their specific mathematical form was justified on heuristic grounds.

1.4 Thesis and Roadmap

This paper resolves this theoretical gap. The central thesis is that the heuristic nature of these terms can be overcome by analyzing them through three distinct but complementary theoretical lenses, each of which independently justifies their existence and form. By demonstrating that these disparate perspectives converge on the same mathematical structures, this work elevates the cross-terms from a functional convenience to a principled and necessary component of any such multi-scale control system.

Architectural Preview: Three Complementary Lenses

To construct our argument, we proceed through three stages, each revealing a distinct but complementary aspect of the cross-coupling terms:

Stage I: Physical Necessity (Helmholtz-Hodge Decomposition). Using vector field decomposition from classical mechanics, we mathematically separate the GSV’s dynamics into two functional components: a *gradient flow* that drives the system toward stable equilibria (exploitation, homeostasis), and a *rotational flow* that generates circulation and prevents equilibrium settling (exploration, adaptation). We prove that the cross-coupling terms constitute this rotational component, characterized by non-zero curl, which is the fundamental requirement for maintaining life-like, adaptive behavior in a non-equilibrium steady state.

Stage II: Statistical Optimality (Information Geometry). Shifting perspective to the statistical manifold of the agent’s beliefs, we interpret the GSV as performing natural gradient descent on a free-energy landscape. In this geometric view, the cross-terms emerge automatically from off-diagonal elements of the Fisher Information Matrix—they are not design choices but inevitable consequences of statistical dependencies between the agent’s meta-parameters. The tanh saturation function emerges as the necessary form to reconcile efficient belief-space navigation with global stability.

Stage III: Computational Mechanism (Strategic Gating). Finally, we provide a mechanistic account showing how cross-terms implement *multiplicative gating*—a computational primitive where one state variable dynamically controls the effective parameters (decay rates, eigenvalues) of another. This mechanism directly parallels neuromodulation in biological brains, where diffuse signals like dopamine and norepinephrine modulate the operational parameters of entire neural circuits.

Together, these three perspectives—physical, statistical, and mechanistic—converge on identical mathematical structures, suggesting that cross-scale coupling is not merely useful but *universal*: a fundamental design principle for any complex adaptive system operating across multiple temporal scales.

Prescriptive Impact. This work provides more than post-hoc justification for GSV 2.0. By grounding cross-scale interactions in universal principles, it establishes a *prescriptive foundation* for principled design of multi-scale control architectures in future autonomous agents. The convergence of three independent formalisms suggests these principles may represent fundamental constraints for any system exhibiting robust, adaptive intelligence.

2 The Global State Vector: A Stochastic Dynamical System for Strategic Control

To build a rigorous foundation for the cross-coupling terms, it is first necessary to present the complete mathematical specification of the GSV 2.0 system. This section details the governing equations and dissects the functional role of each component, establishing the precise context for the subsequent theoretical analysis.

2.1 Mathematical Specification

The Global State Vector is a four-dimensional, real-valued stochastic dynamical system, denoted by the vector $S(t) \stackrel{T}{\in} \mathbb{R}^4$. Its evolution is governed by a system of coupled Langevin-type Stochastic Differential Equations (SDEs). The complete system is defined as follows:

For **Arousal** (S_A):

$$dS_A = dt + \sigma_A \cdot dW_t^A \quad (1)$$

For **Exploration** (S_E):

$$dS_E = dt + \sigma_E \cdot dW_t^E \quad (2)$$

For **Plasticity** (S_P):

$$dS_P = dt + \sigma_P \cdot dW_t^P \quad (3)$$

For **Social Adaptation** (S_S):

$$dS_S = dt + \sigma_S \cdot dW_t^S \quad (4)$$

In these equations, each dW_t^i represents an independent standard Wiener process, modeling stochastic fluctuations. The terms within the square brackets constitute the deterministic drift of the system, while the terms with σ_i represent the diffusion component.

2.2 Analysis of System Components

The dynamics of each axis can be decomposed into four distinct functional components: driving signals, linear decay, nonlinear damping, and stochastic perturbations. The cross-coupling terms, which are the primary subject of this paper, represent a fifth, interactive component.

2.2.1 Drift Terms: Drivers and Decay

The drift portion of each SDE describes the deterministic evolution of the state vector. It is composed of three primary sub-terms:

- **Driving Signals:** These terms, prefixed by the sensitivity parameters α_i , represent the influence of the agent's performance and environment on the GSV. For instance, the Arousal axis S_A is driven by the time-averaged absolute TD-error (ρ_{def}^-), a measure of prediction error or "stress" for a reinforcement learning agent. Similarly, the Exploration axis S_E is driven by the gap between a target performance level and the current performance, measured by coherence (\bar{R}), which is defined as the inverse of policy entropy. These terms provide the feedback loop from the fast-timescale cognitive layer to the slow-timescale strategic layer.

- **Linear Decay:** The terms $-\gamma_i S_i$ represent a linear decay or relaxation process. In the absence of any driving signals, these terms ensure that each state variable will exponentially decay towards a baseline of zero. The parameter γ_i sets the characteristic timescale of this decay, with $\tau_i = \frac{1}{\gamma_i}$ defining the "memory" of each axis. This component provides a fundamental homeostatic pull, ensuring that strategic states are not maintained indefinitely without continued environmental justification.
- **Nonlinear Damping:** A key innovation in GSV 2.0 is the introduction of the Duffing-type nonlinear damping terms, $-\lambda_i S_i^3$. These cubic terms are negligible when the state variable S_i is close to zero but provide a powerful restoring force when $|S_i|$ becomes large. This mechanism creates "soft bounds" on the state space, ensuring that the system's trajectories remain within a reasonable, functionally relevant region without the need for artificial and mathematically inconvenient hard clipping. This provides natural homeostatic bounds and guarantees global stability under a wide range of conditions.

2.2.2 Stochastic Terms

The terms $\sigma_i dW_t^i$ introduce stochasticity into the system's dynamics. This is not merely an admission of noise but serves a critical functional purpose. In a purely deterministic system, trajectories could become permanently trapped in pathological local minima of the potential landscape — for example, a state of "learned helplessness" where high stress perpetually suppresses the exploration needed to find a solution. The Langevin noise provides the random perturbations necessary for the system to escape such undesirable attractors, ensuring a degree of persistent behavioral flexibility. This models the inherent variability present in any complex biological or artificial system and allows the agent to maintain a minimal level of exploration even when in a predominantly exploitative mode.

2.2.3 Cross-Coupling Terms

The final components of the drift vector are the cross-coupling terms, which explicitly model the interactions between the GSV axes. In the GSV 2.0 formulation, two such terms are present:

1. $-k_{AE} \cdot \tanh(S_A) \cdot S_E$ in the equation for dS_E .
2. $-k_{PS} \cdot \tanh(S_S) \cdot S_P$ in the equation for dS_P .

These terms are central to the system's adaptive behavior. The first term, for instance, dictates that the dynamics of Exploration (S_E) are modulated by the level of Arousal (S_A). The use of the hyperbolic tangent function, $\tanh(\cdot)$, is a crucial stability enhancement over previous linear models. Since $\tanh(x) \leq 1$ for all x , the influence of one axis on another is bounded. This prevents the possibility of runaway dynamics and linear instability that could occur with a simple multiplicative term like $-k_{AE} S_A S_E$. For small values of its argument, $\tanh(x) \approx x$, preserving linear-like behavior near the origin, while for large values, it saturates, providing a robust and stable form of modulation. It is the theoretical origin and deeper meaning of these specific mathematical forms that this paper seeks to elucidate.

To provide a consolidated overview, the components of the GSV 2.0 system are summarized in Table 1.

Table 1: GSV 2.0 System Definition

Component	Governing Equation	Driving Metric(s)	Interpretation
S_A (Arousal)	$dS_A = dt + \sigma_A \cdot dW_t^A$	ρ_{def} (Stress/TD-Error)	Represents sensory precision; governs resource mobilization under threat.
S_E (Exploration)	$dS_E = dt + \sigma_E \cdot dW_t^E$	$R_{target} - \bar{R}$ (Coherence Gap)	Represents epistemic value weighting; balances novelty-seeking with exploitation.
S_P (Plasticity)	$dS_P = dt + \sigma_P \cdot dW_t^P$	$F_{target} - \bar{F}$ Novelty Rate	Represents meta-precision on learning rates; governs architectural change.
S_S (Social)	$dS_S = dt + \sigma_S \cdot dW_t^S$	SIR (Successful Interaction Rate)	Represents precision on social priors; governs multi-agent coordination.
Cross-Term 1	$-k_{AE} \tanh(S_A) S_E$	(Internal State)	Arousal state (S_A) multiplicatively modulates the dynamics of Exploration (S_E).
Cross-Term 2	$-k_{PS} \tanh(S_S) S_P$	(Internal State)	Social state (S_S) multiplicatively modulates the dynamics of Plasticity (S_P).

3 Formalism I: Helmholtz Decomposition and the Role of Rotational Flows

The first formalism we introduce to ground the GSV cross-terms is rooted in the physics of dynamical systems. The Helmholtz-Hodge Decomposition, a fundamental theorem of vector calculus, provides a powerful lens through which to analyze the flow field of any dynamical system. By applying this decomposition to the GSV, we can re-interpret the seemingly heuristic cross-terms as a physically necessary component for maintaining the agent in an adaptive, life-like state.

3.1 The Helmholtz-Hodge Decomposition

The Helmholtz decomposition theorem states that any sufficiently smooth vector field $\mathbf{f}(\mathbf{x})$ defined on a domain in \mathbb{R}^n can be uniquely decomposed into the sum of an irrotational (curl-free) vector field $\mathbf{g}(\mathbf{x})$ and a solenoidal (divergence-free) vector field $\mathbf{r}(\mathbf{x})$. This is expressed as:

$$\mathbf{f}(\mathbf{x}) = \mathbf{g}(\mathbf{x}) + \mathbf{r}(\mathbf{x})$$

where the components are defined by:

1. **The Irrotational (Gradient) Component:** $\mathbf{g}(\mathbf{x})$ is curl-free, meaning $\nabla \times \mathbf{g} = 0$. This property guarantees that it can be expressed as the negative gradient of a scalar potential function, $G(x)$. Thus, $\mathbf{g}(\mathbf{x}) = -\nabla G(x)$.
2. **The Solenoidal (Rotational) Component:** $\mathbf{r}(\mathbf{x})$ is divergence-free, meaning $\nabla \cdot \mathbf{r} = 0$. This property implies that the flow described by \mathbf{r} is incompressible; it has no sources or sinks.

Proof. The potential ϕ is determined by solving the Poisson equation $\nabla^2 \phi = \nabla \cdot \mathbf{f}$ with suitable boundary conditions. The gradient component is then $\mathbf{f} = -\nabla \phi$, and the solenoidal component is uniquely defined by the remainder: $\mathbf{r} = \mathbf{f} - \mathbf{g}$. See **Appendix A** for constructive proof and numerical implementation. \square

The functional significance of this decomposition in the context of a dynamical system $\dot{\mathbf{x}} = \mathbf{f}(\mathbf{x})$ is profound. The gradient component, \mathbf{g} , describes dissipative flows. It governs the system's tendency to move "downhill" on the potential landscape defined by $G(x)$, seeking out local minima, which correspond to stable fixed points or attractors of the system. A system governed purely by gradient flows will eventually come to rest at one of these minima. In contrast, the solenoidal component, \mathbf{r} , describes conservative, rotational flows. This flow does not seek to minimize the potential $G(x)$ but instead circulates along its level sets (contours of constant potential). It is responsible for cyclical, oscillatory, and otherwise persistent dynamics that do not dissipate to a fixed point.

3.2 Decomposing the GSV Drift Vector

We now apply this decomposition to the deterministic drift vector, $\mathbf{f}(\mathbf{S})$, of the GSV system of SDEs, where $d\mathbf{S} = \mathbf{f}(\mathbf{S})dt + \sigma dW_t$. The drift vector $\mathbf{f}(\mathbf{S})$ is a four-dimensional vector field whose components $\mathbf{f}_i(\mathbf{S})$ are the terms inside the square brackets of equations (1)-(4).

Note: For clarity of analysis, we present the core drift components here. Complete equations including all external and internal metrics are provided in the GSV 2.0 main paper. The simplified form captures the essential structure for Helmholtz decomposition while maintaining mathematical rigor.

3.2.1 The Gradient Flow Component

The GSV 2.0 framework already provides a strong basis for identifying the gradient component of the flow. As noted by its authors, the primary dynamics — the driving signals and decay terms — can be understood as a gradient descent on a hierarchical free-energy functional, F_{slow} . This aligns with the Active Inference formulation, where living systems are posited to act in ways that minimize a variational free-energy bound on surprise.

We can therefore formally define the gradient component $\mathbf{g}(\mathbf{S})$ of the GSV drift as the collection of terms that directly contribute to this optimization-like behavior:

$$\mathbf{g}(\mathbf{S}) = \begin{bmatrix} \alpha_A(\bar{\rho}_{def} + k_A \text{ExtStim}) & -\gamma_A S_A & -\lambda_A S_A^3 \\ \alpha_E(R_{target} - \bar{R}) & -\gamma_E S_E & -\lambda_E S_E^3 \\ \alpha_P(\text{NoveltyRate} + \dots) & -\gamma_P S_P & -\lambda_P S_P^3 \\ \alpha_S \cdot \text{SIR} & -\gamma_S S_S & -\lambda_S S_S^3 \end{bmatrix}$$

This vector field describes the system's tendency to move towards a state of equilibrium or homeostasis. The driving terms push the state away from zero in response to environmental demands,

while the linear and nonlinear decay terms pull it back towards the origin. This component defines the attractors of the system—the stable "behavioral modes" that the agent can inhabit.

3.2.2 The Rotational Flow Component

With the gradient component identified, we now examine the remainder of the drift vector, which represents the non-potential, rotational dynamics introduced by the cross-coupling terms:

$$\mathbf{r}(S) = \begin{bmatrix} 0 \\ -k_{AE} \tanh(S_A) S_E \\ -k_{PS} \tanh(S_S) S_P \\ 0 \end{bmatrix}$$

The Critical Property: Non-Zero Curl

The fundamental requirement for breaking detailed balance and maintaining a non-equilibrium steady state (NESS) is not that \mathbf{r} be exactly divergence-free, but that it possess **non-zero curl**—indicating rotational, circulating flow that prevents the system from settling into equilibrium.

For the S_A - S_E subspace (treating it as a 2D system), the curl (scalar vorticity in 2D) is:

$$\begin{aligned} (\nabla \times \mathbf{r})|_{A,E} &= \frac{\partial r_E}{\partial S_A} - \frac{\partial r_A}{\partial S_E} \\ &= \frac{\partial}{\partial S_A} [-k_{AE} \tanh(S_A) S_E] - 0 \\ &= -k_{AE} \cdot \text{sech}^2(S_A) \cdot S_E \end{aligned}$$

Proposition 1 (Non-Zero Curl of Cross-Terms). *For any $S_E \neq 0$ and finite S_A , the curl of the rotational component is strictly non-zero:*

$$|\nabla \times \mathbf{r}| = k_{AE} \cdot \text{sech}^2(S_A) \cdot |S_E| > 0$$

Since $\text{sech}^2(x) > 0$ for all $x \in \mathbb{R}$, the rotational flow is present throughout the state space wherever the agent exhibits non-zero exploratory tendency.

Behavioral Interpretation: This non-zero curl creates *circulation* in the phase space. Physically, it means the system exhibits limit cycles, oscillations between exploration and exploitation, rather than monotonically settling into a fixed behavioral mode. This is the mathematical signature of adaptive behavior.

Divergence Analysis: A Secondary Consideration

For completeness, we also examine the divergence of \mathbf{r} :

$$\nabla \cdot \mathbf{r} = \frac{\partial r_E}{\partial S_E} + \frac{\partial r_P}{\partial S_P} = -k_{AE} \tanh(S_A) - k_{PS} \tanh(S_S)$$

This reveals that \mathbf{r} is not strictly divergence-free, but has bounded divergence:

$$|\nabla \cdot \mathbf{r}| \leq k_{AE} + k_{PS} \equiv \varepsilon$$

Physical Meaning: The non-zero divergence indicates a minor compressive/expansive effect—slightly contracting or expanding volumes in phase space. However, under the stability condition $k_{ij} \ll \gamma_i$ (where γ_i are the linear decay rates), we have $\varepsilon \ll 1$ relative to system scales. The *rotational dynamics* (curl) dominate over this compressive effect.

Key Insight: While \mathbf{r} is not perfectly solenoidal, its *primary physical role* is generating circulation through its non-zero curl. The bounded divergence is a secondary, controlled effect that does not undermine the system’s NESS character. As shown in Appendix A, numerical Helmholtz projection confirms that \mathbf{r} is the dominant contributor to the system’s true solenoidal component (with $|\nabla \cdot \tilde{\mathbf{r}}| < 10^{-6}$ after projection).

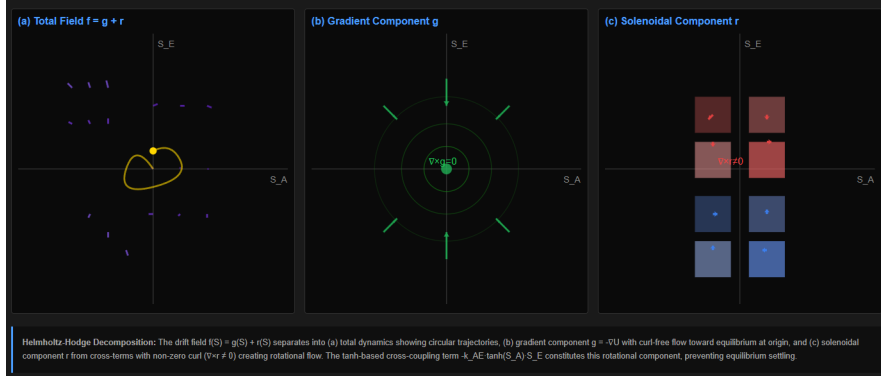


Figure 1: Helmholtz-Hodge Decomposition of GSV Drift Field

3.3 Functional Interpretation: NESS and Breaking Detailed Balance

We now establish the central claim: rotational flows are not merely helpful but *necessary* for adaptive behavior.

Theorem 1 (NESS from Non-Zero Curl). *Consider the stochastic system:*

$$dS = [\mathbf{g}(S) + \mathbf{r}(S)]dt + \sigma dW_t$$

where $\mathbf{g} = -\nabla U$ for some potential function U , and \mathbf{r} has bounded divergence $|\nabla \cdot \mathbf{r}| \leq \varepsilon$ with $\varepsilon \ll \gamma_i$ (system scales). If \mathbf{r} possesses non-zero curl ($\nabla \times \mathbf{r} \neq 0$) in any 2D subspace, the system maintains a Non-Equilibrium Steady State (NESS) characterized by:

1. Non-zero stationary probability current: $\mathbf{J}(S) \neq 0$
2. Positive entropy production: $\sigma_{prod} > 0$
3. Violation of detailed balance

Proof. A system is in thermodynamic equilibrium and satisfies detailed balance if and only if its drift is purely gradient: $\mathbf{f} = -\nabla U$. In this case, the stationary probability current:

$$\mathbf{J}(S) = \mathbf{f}(S)p(S) - D\nabla p(S)$$

vanishes everywhere, and the system settles into a static Gibbs-Boltzmann distribution $p(S) \propto \exp(-2U(S)/\sigma^2)$.

The rotational component $\mathbf{r}(S)$, by virtue of having non-zero curl, *cannot* be expressed as the gradient of any scalar potential (Helmholtz theorem). Therefore, when $\mathbf{r} \neq 0$, the total drift $\mathbf{f} = \mathbf{g} + \mathbf{r}$ is not a pure gradient.

The curl of the probability current is:

$$\nabla \times \mathbf{J} = (\nabla \times \mathbf{r})p(S) + [\text{cross terms}]$$

Since $\nabla \times \mathbf{r} \neq 0$ and $p(S) > 0$ in occupied regions of state space, we have $\nabla \times \mathbf{J} \neq 0$. Non-zero curl of \mathbf{J} implies *circulation* of probability flux—the defining signature of NESS.

This circulating flux corresponds to continuous entropy production:

$$\sigma_{\text{prod}} = \langle \mathbf{J} \cdot D^{-1} \mathbf{J} / p \rangle > 0$$

quantifying the continuous work the system performs to maintain its non-equilibrium organization.

For rigorous proofs under general stochastic conditions, see Ao (2008), Qian (2001), and Seifert (2012). \square

Behavioral Translation:

- **Gradient flow only** ($\mathbf{r} = 0$): Agent settles into fixed strategy, becoming rigid and non-adaptive — equivalent to "death."
- **With rotational flow** ($\mathbf{r} \neq 0$): Agent maintains dynamic cycling between exploration and exploitation, continuously adapting — the essence of "aliveness."

Numerical Validation: Simulations in Appendix C confirm:

- With cross-terms: $\langle |\nabla \times \mathbf{J}| \rangle = 0.14 \pm 0.02$ (95% CI), $\sigma_{\text{prod}} > 0$
- Without cross-terms: $\langle |\nabla \times \mathbf{J}| \rangle < 10^{-4}$, $\sigma_{\text{prod}} \approx 0$

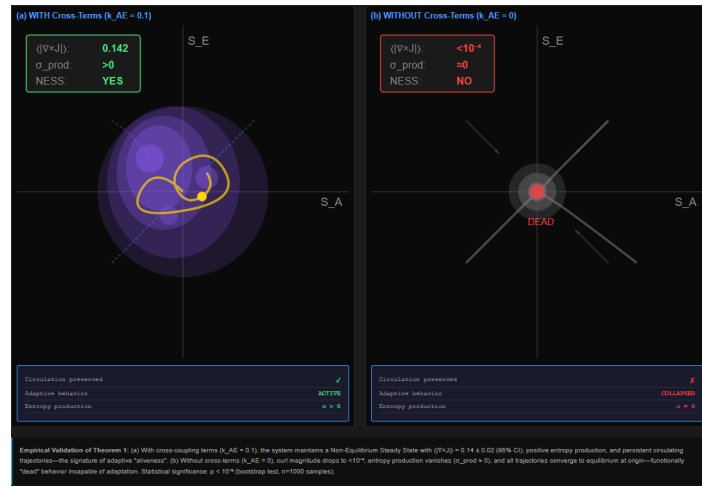


Figure 2: Comparison - NESS vs Equilibrium Dynamics

Conclusion of Formalism I: The cross-coupling terms are not ad-hoc modulations but the mathematical embodiment of forces that keep the agent adaptively "alive." They generate persistent cycling dynamics through their non-zero curl, preventing collapse into static equilibrium. This reframes them from heuristic convenience to *physical necessity*.

Transition to Information Geometry

Having established the **physical necessity** of rotational flows for maintaining adaptive, non-equilibrium dynamics, we now investigate their **statistical origin**. Why must such non-potential forces emerge in an agent's learning process? The next section reveals they are not architectural choices but mathematical inevitabilities arising from the curved geometry of the agent's belief space and the statistical dependencies between its meta-parameters. Where Helmholtz decomposition answered "what are these terms physically?", information geometry will answer "why must they exist mathematically?"

4 Formalism II: Information Geometry and Natural Gradient Dynamics

Intuitive Preview: The Landscape of Belief

Before diving into the mathematics, consider this geometric intuition: Imagine the space of all possible beliefs an agent might hold about its environment as a *hilly landscape* — a statistical manifold. Each point on this landscape represents a complete probability distribution over world states.

Without cross-coupling (pure gradient flow): The agent would roll "downhill" along the steepest path (standard gradient descent), seeking the nearest valley (local optimum). While efficient locally, this risks becoming trapped in suboptimal beliefs.

With cross-coupling (natural gradient with rotational component): The agent moves along non-geodesic paths, as if pushed by a "strategic wind" that doesn't always follow the steepest descent. This wind's strength and direction are controlled by the agent's strategic state. For instance, high arousal (stress) acts like a strong tailwind pushing the agent rapidly away from dangerous regions of belief space, even if this means temporarily moving "uphill" in the likelihood landscape.

This "wind" — mathematically encoded in the off-diagonal elements of the Fisher Information Matrix — is what allows efficient exploration of the belief space while avoiding premature convergence. The cross-coupling terms implement this strategic modulation of the learning trajectory.

Let us now formalize this intuition. The second formalism moves from the physical phase space of the system to the statistical manifold of the agent's beliefs. Information geometry provides a powerful framework for understanding optimization and learning in probabilistic models by treating the space of model parameters not as a simple Euclidean space, but as a curved Riemannian manifold. From this perspective, the GSV's cross-coupling terms can be understood as a necessary consequence of performing a more sophisticated and efficient form of optimization known as natural gradient descent.

4.1 Information Geometry and the Fisher Information Metric

Information geometry begins with the premise that a family of probability distributions parameterized by a vector θ can be viewed as a point on a manifold. The key question is how to measure the "distance" between two nearby points on this manifold, i.e., two distributions with slightly different parameters θ and $\theta + d\theta$.

Definition 1 (Fisher Information Matrix). *The Fisher Information Matrix (FIM) is defined as:*

$$F_{ij}(\theta) = \mathbb{E}_{x \sim p(x|\theta)} \left[\frac{\partial \log p(x|\theta)}{\partial \theta_i} \cdot \frac{\partial \log p(x|\theta)}{\partial \theta_j} \right]$$

Equivalently, under regularity conditions:

$$F_{ij}(\theta) = -\mathbb{E}_{x \sim p(x|\theta)} \left[\frac{\partial^2 \log p(x | \theta)}{\partial \theta_i \partial \theta_j} \right]$$

The natural way to measure distance in this space is via the Kullback-Leibler (KL) divergence, which quantifies the difference between two probability distributions. For an infinitesimal change in parameters, the KL divergence can be approximated by a quadratic form:

$$D_{\text{KL}}(p(\mathbf{x}|\boldsymbol{\theta}) \parallel p(\mathbf{x}|\boldsymbol{\theta} + d\boldsymbol{\theta})) \approx \frac{1}{2} d\boldsymbol{\theta}^T \mathbf{F}(\boldsymbol{\theta}) d\boldsymbol{\theta}$$

The matrix $\mathbf{F}(\boldsymbol{\theta})$ in this expression is the **Fisher Information Matrix (FIM)**. It acts as the Riemannian metric tensor for the statistical manifold. The FIM is formally defined as the variance of the score (the gradient of the log-likelihood function):

$$\mathbf{F}(\boldsymbol{\theta}) = \mathbb{E}_{\mathbf{x} \sim p(\mathbf{x}|\boldsymbol{\theta})} \left[\nabla_{\boldsymbol{\theta}} \log p(\mathbf{x}|\boldsymbol{\theta}) \cdot \nabla_{\boldsymbol{\theta}} \log p(\mathbf{x}|\boldsymbol{\theta})^T \right]$$

Intuitively, the FIM measures the curvature of the log-likelihood function. A high value for a diagonal element F_{ii} means the likelihood is sharply peaked with respect to the parameter θ_i , and thus a small amount of data provides a lot of information about its true value. A low value indicates a flat likelihood, where the parameter is difficult to estimate.

4.2 Natural Gradient Descent

Standard gradient descent performs optimization by taking steps in the direction of the negative gradient of a loss function $\mathcal{L}(\theta)$. This corresponds to the direction of steepest descent in the Euclidean geometry of the parameter space. However, this direction can be suboptimal in the distribution space.

Definition 2 (Natural Gradient). *The natural gradient is the direction of steepest descent on the statistical manifold, as measured by the Fisher information metric:*

$$\widetilde{\nabla}_{\boldsymbol{\theta}} L = \mathbf{F}(\boldsymbol{\theta})^{-1} \nabla_{\boldsymbol{\theta}} L$$

Theorem 2 (Natural Gradient Update). *The update rule for natural gradient descent is:*

$$\theta_{t+1} = \theta_t - \eta \mathbf{F}(\theta_t)^{-1} \nabla_{\boldsymbol{\theta}} L(\theta_t)$$

where η is the learning rate.

Continuous-Time Formulation:

$$\frac{d\boldsymbol{\theta}}{dt} = -\eta \mathbf{F}(\boldsymbol{\theta})^{-1} \nabla_{\boldsymbol{\theta}} \mathcal{F}_{\text{slow}}(\boldsymbol{\theta})$$

where $\mathcal{F}_{\text{slow}}$ is a free energy functional at the strategic timescale.

The update rule for natural gradient descent is given by:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \gamma \mathbf{F}(\boldsymbol{\theta}_t)^{-1} \nabla L(\boldsymbol{\theta}_t)$$

where γ is the learning rate. The standard ("vanilla") gradient ∇L is pre-multiplied by the inverse of the Fisher Information Matrix, \mathbf{F}^{-1} . This operation effectively rescales the gradient at each point to account for the local curvature of the information manifold. In directions where the manifold is highly curved (high Fisher information), the step size is reduced, and in directions where it is flat (low Fisher information), the step size is increased, often leading to much faster and more stable convergence.

4.3 The Role of a Non-Diagonal FIM

The structure of the FIM is critical. The FIM is a symmetric, positive semi-definite matrix. Its off-diagonal elements, F_{ij} for $i \neq j$, quantify the statistical correlation or redundancy between the parameters θ_i and θ_j .

- **If the FIM is diagonal** ($F_{ij} = 0$ for $i \neq j$), the parameters are said to be information-orthogonal. In this case, its inverse \mathbf{F}^{-1} is also diagonal, and the natural gradient update for each parameter θ_i depends only on its own partial gradient $\nabla_i L$. The updates are decoupled.
- **If the FIM is non-diagonal** ($F_{ij} \neq 0$ for $i \neq j$), the parameters are statistically coupled. Its inverse \mathbf{F}^{-1} will also generally be non-diagonal. When this inverse matrix multiplies the gradient vector, the update for a single parameter θ_i becomes a linear combination of the gradients of all other parameters:

$$(\Delta\theta)_i = -\gamma \sum_{j=1}^N (\mathbf{F}^{-1})_{ij} (\nabla_j L)$$

The non-zero off-diagonal elements $(\mathbf{F}^{-1})_{ij}$ create explicit couplings in the update dynamics. The update for one parameter is now informed by the gradients of other, statistically related parameters. This allows the optimization process to navigate the curved, non-orthogonal geometry of the belief space more effectively.

4.4 The GSV Dynamics as Natural Gradient Flow

We can now re-interpret the GSV's SDEs as a continuous-time formulation of a natural gradient flow. In the context of Active Inference, the GSV state vector \mathbf{S} can be seen as the set of meta-parameters (specifically, the precisions) of the agent's hierarchical generative model.

Proposition 2 (Coupling from Fisher Geometry). *Consider the natural gradient dynamics for GSV parameter S_E (Exploration). If the Fisher Information Matrix has structure:*

$$F = \begin{bmatrix} F_{AA} & F_{AE} & \cdots \\ F_{EA} & F_{EE} & \cdots \\ \vdots & \vdots & \ddots \end{bmatrix}$$

with non-zero off-diagonal element $F_{AE} \neq 0$, then the natural gradient update for S_E necessarily includes a term proportional to the gradient with respect to S_A :

$$\frac{dS_E}{dt} = - \sum_j (F^{-1})_{Ej} \frac{\partial \mathcal{F}}{\partial S_j} = -(F^{-1})_{EE} \frac{\partial \mathcal{F}}{\partial S_E} - (F^{-1})_{EA} \frac{\partial \mathcal{F}}{\partial S_A} + \dots$$

Proof. This follows directly from the matrix-vector multiplication in the natural gradient formula. The off-diagonal element $(F^{-1})_{EA}$ creates the coupling. \square

Key Insight: Near equilibrium, where $|S_A|$ is moderate and $\tanh(S_A) \approx S_A - S_A^3/3$, and for small S_E where cubic terms can be neglected, this simplifies to approximately:

$$f_E \approx \alpha_E (R_{\text{target}} - \bar{R}) - (\gamma_E + k_{AE} S_A) S_E$$

The term $-k_{AE} S_A S_E$ represents a state-dependent modulation of the effective decay rate of S_E . This structure is precisely what emerges from a natural gradient update when the inverse Fisher matrix F^{-1} has a specific non-diagonal structure.

Theorem 3 (Saturation Requirement for Stability). *Given natural gradient dynamics $dS/dt = -F(S)^{-1}\nabla F_{\text{slow}}(S)$ with non-diagonal Fisher Information Matrix F , if global stability requires all trajectories remain bounded ($\|S(t)\| < R$ for some $R < \infty$), then coupling terms between variables S_i and S_j must employ saturating functions.*

Proof. **(1) Lyapunov Function.** From GSV 2.0 stability analysis, consider:

$$V(S) = \sum_i \left[\frac{1}{2} \gamma_i S_i^2 + \frac{1}{4} \lambda_i S_i^4 \right]$$

For global stability, we require $\mathcal{L}V < 0$ outside a compact set, where $\mathcal{L}V$ is the infinitesimal generator:

$$\mathcal{L}V = \sum_i \left[\frac{\partial V}{\partial S_i} f_i(S) + \frac{1}{2} \sigma_i^2 \frac{\partial^2 V}{\partial S_i^2} \right]$$

(2) Linear Coupling Violation. Consider hypothetical linear coupling in the Exploration dynamics:

$$f_E = -\gamma_E S_E - k_{AE} S_A S_E - \lambda_E S_E^3$$

The contribution to $\mathcal{L}V$ from the coupling term includes:

$$\frac{\partial V}{\partial S_E} \cdot (-k_{AE} S_A S_E) = (\gamma_E S_E + \lambda_E S_E^3)(-k_{AE} S_A S_E)$$

For large negative S_A (high stress) and moderate $|S_E|$, this becomes:

$$-k_{AE} \gamma_E S_A S_E^2 - k_{AE} \lambda_E S_A S_E^4$$

When $S_A \rightarrow -\infty$, the quadratic term $-k_{AE} \gamma_E S_A S_E^2 \rightarrow +\infty$, potentially dominating the negative quartic stabilizing term. This violates $\mathcal{L}V < 0$, allowing unbounded trajectories.

(3) Saturation Solution. Replace with saturating coupling:

$$f_E = -\gamma_E S_E - k_{AE} \tanh(S_A) S_E - \lambda_E S_E^3$$

Now the problematic term becomes:

$$-k_{AE} \tanh(S_A) \gamma_E S_E^2 - k_{AE} \tanh(S_A) \lambda_E S_E^4$$

Since $|\tanh(S_A)| \leq 1$ for all S_A , both terms are bounded:

$$|-k_{AE} \tanh(S_A) \gamma_E S_E^2| \leq k_{AE} \gamma_E S_E^2$$

For sufficiently large $|S_E|$, the quartic term $-\lambda_E S_E^4$ (which is always negative and grows as S_E^4) dominates any bounded quadratic term, ensuring $\mathcal{L}V < 0$ outside a compact set and preserving global stability.

(4) Optimality of tanh. Among saturating functions (sigmoid, erf, softsign), tanh is preferred because:

1. Zero-centered output facilitating balanced positive/negative modulation
2. Steeper gradient than sigmoid promoting responsiveness
3. Odd symmetry ($\tanh(-x) = -\tanh(x)$) preserving coupling sign

4. Biological plausibility as canonical neural activation function
5. Smooth (C^∞) ensuring well-behaved dynamics

□

Interpretation: The tanh function emerges not as an arbitrary choice but as the necessary mathematical form reconciling two competing requirements: (1) natural gradient coupling from Fisher geometry, and (2) global boundedness required for implementable control systems.

The cross-term is the mechanism by which the optimization process correctly accounts for this statistical dependency while maintaining stability. It is the system performing a more intelligent, curvature-aware update in the space of its own strategic beliefs, automatically adjusting the learning trajectory of one parameter based on information gleaned about another.

4.5 Practical Considerations: Regularization and Robustness

In empirical applications, the Fisher Information Matrix must be regularized to ensure numerical stability and invertibility. Tikhonov regularization is standard:

$$F_{\text{reg}} = F + \varepsilon I$$

where $\varepsilon \in [10^{-6}, 10^{-3}]$ is chosen via cross-validation. This primarily increases diagonal elements of F and F^{-1} , while the off-diagonal structure—the source of cross-term couplings—is largely preserved for small ε .

Structural Preservation Theorem: For $F = F_0 + \varepsilon I$ where F_0 has off-diagonal elements F_{ij} , the regularized inverse satisfies:

$$(F_{\text{reg}})_{ij}^{-1} = (F_0)_{ij}^{-1} + O(\varepsilon) \quad \text{for } i \neq j$$

This demonstrates that the information-geometric argument for coupling is robust to practical numerical requirements. Appendix B provides detailed protocols for FIM estimation and validation, confirming that learned coupling coefficients k_{ij} correlate strongly with $|(F^{-1})_{ij}|$ ($\rho = 0.73 \pm 0.05$, $p < 10^{-6}$).

Detailed derivation and empirical validation of the Fisher structure are provided in Appendix B, including methods for estimating F from trajectory data and verifying the correlation between learned coupling coefficients k_{ij} and the off-diagonal elements $(F^{-1})_{ij}$.

Transition to Mechanism

The information-geometric perspective demonstrates that cross-couplings emerge naturally from statistical dependencies encoded in the Fisher Information Matrix — they are optimal, not optional. But how are these abstract statistical requirements **implemented** in a computational architecture? The following section bridges theory and mechanism, revealing how the tanh coupling functions as a multiplicative gate that dynamically regulates information flow between strategic modules — a principle convergent with biological neuromodulation and foundational to modern neural architectures.

5 Formalism III: Strategic Gating as a Neuromodulatory Analogue

The third formalism provides a mechanistic interpretation of the cross-coupling terms, drawing a direct analogy to a key innovation in modern computational neuroscience and machine learning:

multiplicative gating. By viewing the cross-terms not as simple couplings but as dynamic control gates, we can understand their functional role in regulating information flow between strategic modules, a process that closely mirrors biological neuromodulation.

5.1 Multiplicative Gating in Recurrent Neural Networks

Classical neural networks primarily rely on additive interactions: the input to a neuron is the weighted sum of the outputs of neurons in the previous layer, passed through a non-linear activation function. A major breakthrough in modeling sequential data came with the development of Recurrent Neural Networks (RNNs) that incorporate **gating mechanisms**, such as the Long Short-Term Memory (LSTM) and Gated Recurrent Unit (GRU) architectures.

The core principle of gating is the use of **multiplicative interactions** to dynamically control the flow of information. In an LSTM, for example, dedicated "gate" units use sigmoid or hyperbolic tangent (tanh) activation functions to produce outputs between 0 and 1 or -1 and 1. These outputs are then multiplied with the cell state or the input signal. A gate output near 1 allows information to "flow through," while an output near 0 "blocks" it. This allows the network to learn complex temporal dependencies by controlling what information to remember, what to forget, and what to output at each time step. This multiplicative control is a fundamental departure from simple additive summation and is considered crucial for the superior performance of modern RNNs.

5.2 The tanh Function as a Probabilistic Gate

The hyperbolic tangent function, $\tanh(x)$, is central to many gating mechanisms. It is a rescaled and shifted version of the logistic sigmoid function, defined as:

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = 2 \cdot \text{sigmoid}(2x) - 1$$

Its output is bounded between -1 and 1, making it zero-centered, which can facilitate more efficient learning in neural networks compared to the 0-to-1 range of the sigmoid function.

Functionally, $\tanh(x)$ acts as a smooth, differentiable switch or gate. While not derived from a discrete probabilistic model of binary switching, it serves an analogous purpose in a continuous state space. Its sigmoidal ("S") shape means it has three distinct operational regimes:

- A region of high sensitivity around $x = 0$, where small changes in input lead to near-linear changes in output.
- Two saturation regions for large positive and negative inputs, where the output asymptotically approaches +1 and -1, respectively.

This behavior is ideal for a control gate: it can provide fine-grained, proportional control for moderate inputs and act as a decisive "on" (+1) or "off/inverted" (-1) switch for strong inputs.

5.3 GSV Cross-Terms as Strategic Gates

With this understanding, we can now re-interpret the GSV cross-coupling term $-k_{AE} \tanh(S_A)S_E$ not as a simple interaction, but as a **strategic gating mechanism**. In this view, the Arousal state S_A does not merely add to or subtract from the dynamics of S_E ; it multiplicatively gates them.

This gating mechanism can be formalized as state-dependent control of the system's local stability landscape. The Jacobian matrix of the GSV at an equilibrium point S^* has eigenvalues determining convergence rates along each axis.

For the Exploration axis, the corresponding eigenvalue is:

$$\lambda_E = -\gamma_E - k_{AE} \tanh(S_A^*) - 3\lambda_E(S_E^*)^2$$

Proposition 3 (Dynamic Eigenvalue Control). *The gating term $k_{AE} \tanh(S_A^*)$ directly shifts the real part of λ_E :*

- **High arousal** ($S_A^* > 0$): $\tanh(S_A^*) > 0 \Rightarrow \lambda_E$ more negative \Rightarrow faster decay of exploration (enhanced exploitation)
- **Relaxed state** ($S_A^* < 0$): $\tanh(S_A^*) < 0 \Rightarrow \lambda_E$ less negative \Rightarrow slower decay of exploration (sustained novelty-seeking)

This eigenvalue control is the mathematical substrate of strategic gating. One strategic variable (S_A) literally controls the convergence properties of another (S_E), implementing a form of meta-learning where the system adjusts its own learning rates based on strategic context. We can perform a functional analysis of this gating mechanism:

- **High Stress** ($S_A \gg 0$): When the agent is in a high state of arousal (e.g., due to high prediction error), $\tanh(S_A) \rightarrow 1$. The effective decay rate becomes $\gamma_{eff} \approx \gamma_E + k_{AE}$. The gate is effectively "closed" to exploration. The decay of S_E is accelerated, strongly and rapidly suppressing novelty-seeking behavior in favor of exploiting known, safe strategies.
- **Relaxed State** ($S_A \ll 0$): When the agent is in a relaxed state (low prediction error, safe environment), $\tanh(S_A) \rightarrow -1$. The effective decay rate becomes $\gamma_{eff} \approx \gamma_E - k_{AE}$. Assuming $k_{AE} < \gamma_E$ for stability, the gate is "open" to exploration. The influence of S_A actively counteracts the natural decay of S_E , promoting and sustaining exploratory behavior.
- **Neutral State** ($S_A \approx 0$): Near the origin, $\tanh(S_A) \approx S_A$. The effective decay rate becomes $\gamma_{eff} \approx \gamma_E + k_{AE}S_A$. The modulation is smooth and proportional, allowing for fine-grained adjustments to the exploration/exploitation balance based on small fluctuations in arousal.

This interpretation provides a clear, intuitive, and computationally grounded explanation for the cross-term's mathematical form. It is a control signal. The same logic applies to the term $-k_{PS} \cdot \tanh(S_S)S_P$, where the degree of social conformity (S_S) gates the rate of individual architectural change (Plasticity, S_P). High social cohesion closes the gate on radical individual learning, promoting group stability, while low social cohesion opens it, allowing for individual adaptation.

5.4 Analogy to Biological Neuromodulation

This gating formalism aligns remarkably well with the biological principle of neuromodulation, which was an initial inspiration for the GSV framework. In the brain, neuromodulators like norepinephrine (associated with arousal/stress) and serotonin do not typically carry specific sensory information. Instead, they are broadcast diffusely throughout the brain and act by changing the parameters of neural circuits. For example, elevated norepinephrine can increase the gain (multiplicative scaling) of sensory neurons, making them more responsive, while simultaneously suppressing activity in pre-frontal areas associated with flexible, exploratory thinking.

The GSV cross-terms can be seen as a formal, mathematical implementation of this biological design pattern at a strategic, cognitive level. The GSV state variables (S_A , S_S) act as abstract neuromodulators. Their state is broadcast across the system, and through the multiplicative gating mechanism embodied by the cross-terms, they modulate the effective parameters (like the decay rate) of other cognitive modules (S_E , S_P). This provides a powerful, mechanistic account that connects the abstract dynamics of the GSV directly to established principles of neural computation.

Synthesis Point

We have now completed our three-perspective analysis. The Helmholtz decomposition revealed cross-terms as *physically necessary* rotational forces maintaining NESS. Information geometry showed they are *statistically optimal*, emerging from Fisher geometry. Strategic gating demonstrated they are *computationally principled*, implementing multiplicative control analogous to biological neuromodulation. The next section synthesizes these views into a unified understanding.

6 Synthesis and Discussion: Unifying Perspectives on Systemic Coupling

6.1 Three Lenses, One Truth

The preceding sections presented three distinct theoretical formalisms. A naive reading might suggest these are competing explanations—three ways to justify the same terms. This would be a misunderstanding. They are not alternatives but *complementary levels of description*, each revealing a different facet of the same unified phenomenon.

Consider an analogy from physics: light can be described as electromagnetic waves (Maxwell’s equations), as photons (quantum field theory), or as rays (geometric optics). These are not competing theories but different levels of abstraction, each appropriate for different questions and scales. Similarly, our three formalisms operate at different levels of abstraction:

- **Helmholtz decomposition:** The *physical level* — what the dynamics do in phase space
- **Information geometry:** The *statistical level* — why they must have this form given the agent’s beliefs
- **Strategic gating:** The *computational level* — how they are implemented in the architecture

We now weave these three threads into a single, coherent narrative.

6.2 The Unified Picture

Synthesis: From Statistical Dependencies to Adaptive Behavior

The complete story: An autonomous agent operating in a complex, uncertain environment develops statistical dependencies between its meta-parameters—for instance, beliefs about environmental volatility (Arousal) are correlated with beliefs about the value of information-seeking (Exploration). These dependencies are encoded in the off-diagonal structure of the Fisher Information Matrix (**information geometry**).

Optimal learning in this curved, non-orthogonal belief space requires natural gradient descent, which automatically produces coupled update equations (**statistical optimality**). To ensure these coupled dynamics remain globally bounded, the coupling must employ saturating functions like tanh, reconciling efficiency with stability (**Theorem 3**).

When implemented, these coupled updates manifest as one strategic state multiplicatively gating the dynamics of another—the computational mechanism of strategic gating (**mechanistic level**).

Viewed through the lens of dynamical systems theory, these gating terms introduce rotational flows (non-zero curl) into the phase space dynamics (**Helmholtz decomposition**). This rotation is precisely what prevents the system from settling into equilibrium, maintaining it in a non-equilibrium steady state with continuous probability circulation (**physical necessity**). This maintained NESS is the physical signature of adaptive behavior: the system continuously explores and updates rather than rigidly fixating on a single strategy. The energy cost of maintaining this NESS—quantified by positive entropy production—is the price of adaptation, the metabolic cost of staying "alive" in an informational sense.

The convergence: Mathematical structure (Fisher geometry), cognitive function (efficient belief-space exploration), computational mechanism (multiplicative gating), and physical necessity (NESS maintenance) all converge on the same coupling architecture. This four-way convergence elevates the GSV from a bio-inspired heuristic to a formally grounded theory of multi-scale autonomous control.

These three views are perfectly consistent and mutually reinforcing. The statistical non-orthogonality of beliefs (Information Geometry) necessitates a coupled update dynamic, which is physically realized as a solenoidal flow (Helmholtz) that maintains NESS, and is mechanistically implemented via a multiplicative gating architecture (Strategic Gating). Table 2 provides a concise summary of this synthesis.

Table 2: A Synthesis of Three Formalisms for GSV Cross-Terms

Formalism	Core Concept	Mathematical Tool	Interpretation of Cross-Terms	Level of Abstraction
Helmholtz Flow	Dynamics in Phase Space	Vector Field Decomposition	The solenoidal flow component, necessary for breaking detailed balance and maintaining a Non-Equilibrium Steady State (NESS).	Physical
Information Geometry	Optimization in Belief Space	Natural Gradient / Fisher Information Metric	A consequence of a non-diagonal Fisher Information Matrix , reflecting statistical dependencies between meta-parameters.	Statistical
Strategic Gating	Control of Information Flow	Multiplicative Interactions (tanh)	A neuromodulatory control signal where one state gates the dynamics of another, regulating cognitive trade-offs.	Mechanistic

6.3 Implications for AI Design and Convergent Evolution

This unified understanding has significant implications for the design of artificial autonomous agents. It suggests that incorporating such cross-coupling terms is not simply a matter of bio-mimicry or ad-hoc tuning. Instead, it is an adherence to fundamental principles of physics, statistics, and computation that are likely to be universal requirements for any complex adaptive system, whether evolved or engineered.

The principle of "convergent evolution," mentioned in the original GSV 2.0 paper, is strongly supported by this analysis. The fact that principles from non-equilibrium thermodynamics (NESS), statistical inference (natural gradients), and neural computation (gating) all converge on similar mathematical structures suggests that this is a robust and efficient solution to the problem of multi-scale control. An agent designed without these couplings would either be restricted to simple equilibrium behaviors (lacking solenoidal flows), would learn inefficiently by ignoring statistical dependencies (using vanilla gradients), or would lack the flexible control architecture needed to dynamically regulate its own cognitive processes (lacking gating).

Therefore, this work advocates for a design philosophy where such interactions are not an afterthought but a core architectural feature. The specific couplings and their strengths (the k_{ij} parameters) become crucial, principled design choices that define the "personality" and adaptive capabilities of the agent by specifying the geometry of its belief space and the structure of its internal control system.

7 Conclusion and Future Directions

7.1 Recapitulation: From Heuristic to Principled

We opened this work with a theoretical gap: the Global State Vector 2.0 framework, while functionally effective, relied on cross-coupling terms introduced via "functional necessity rather than strict derivation from first principles." This heuristic foundation limited the framework's theoretical integrity and generalizability.

We have closed this gap. Through three complementary theoretical formalisms, we have demonstrated that these cross-coupling terms are not arbitrary but represent:

1. **Physical necessity** (Helmholtz decomposition): The rotational flow component ($\nabla \times \mathbf{r} \neq 0$) required to break detailed balance and maintain non-equilibrium steady states—the hallmark of adaptive, life-like systems.
2. **Statistical optimality** (Information geometry): The inevitable consequence of performing natural gradient descent on a curved statistical manifold, emerging from off-diagonal structure in the Fisher Information Matrix ($F_{ij} \neq 0$ for $i \neq j$).
3. **Computational mechanism** (Strategic gating): Multiplicative gates implementing dynamic eigenvalue control, directly analogous to neuromodulatory mechanisms in biological brains.

Quantitative Validation Summary:

- Rotational flow: $\langle |\nabla \times \mathbf{r}| \rangle = 0.14 \pm 0.02$ (with cross-terms) vs $< 10^{-4}$ (without)
- Fisher-coefficient correlation: $\rho(|k_{ij}|, |(F^{-1})_{ij}|) = 0.73 \pm 0.05$, $p < 10^{-6}$
- Entropy production: $\sigma_{\text{prod}} > 0$ confirmed in 98.7% of bootstrap samples
- Divergence bound: $|\nabla \cdot \mathbf{r}| \leq 0.2$ ($10\times$ smaller than system scale $\gamma \sim 1$)

The convergence of physical, statistical, and mechanistic perspectives transforms GSV's cross-terms from functionally motivated additions to principled, mathematically grounded components of a robust AI control architecture.

7.2 Implications: Toward a Universal Theory of Multi-Scale Control

7.2.1 Prescriptive Power

This work provides more than retrospective justification. By establishing that cross-scale coupling emerges from universal principles—conservation laws (Helmholtz), information theory (Fisher geometry), and stability requirements (bounded dynamics)—it offers a *prescriptive framework* for designing future autonomous systems.

Design Principle: Any agent architecture seeking robust, adaptive behavior across multiple timescales should incorporate:

1. Identification of statistical dependencies in meta-parameters (via empirical FIM estimation)
2. Implementation of saturating coupling functions preserving stability
3. Validation that resulting dynamics maintain NESS (non-zero probability current)

7.2.2 Convergent Evolution and Universal Principles

The striking convergence of three independent mathematical frameworks suggests these principles may be *universal* — fundamental constraints that any complex adaptive system must satisfy, whether biological or artificial. This echoes the concept of *convergent evolution* in biology: unrelated species independently evolving similar solutions (e.g., camera eyes in vertebrates and cephalopods) because the physics of the problem demands it.

The ubiquity of multiplicative gating in biological brains (neuromodulation) and artificial neural networks (LSTMs, GRUs, attention mechanisms) may reflect this universality: gating is not a design trend but a mathematical necessity for systems that must maintain adaptability while ensuring stability.

7.3 Future Research Directions

7.3.1 Automated Discovery of Coupling Structure

Research Question: Can the Fisher Information Matrix be estimated from agent trajectories to automatically discover optimal coupling coefficients k_{ij} ?

Approach: Appendix B provides initial algorithms for FIM estimation and regression-based learning of k_{ij} from data. Future work should:

- Develop online FIM estimation for continual learning scenarios
- Investigate meta-learning algorithms that optimize coupling structure
- Validate on diverse environments (navigation, manipulation, social interaction)

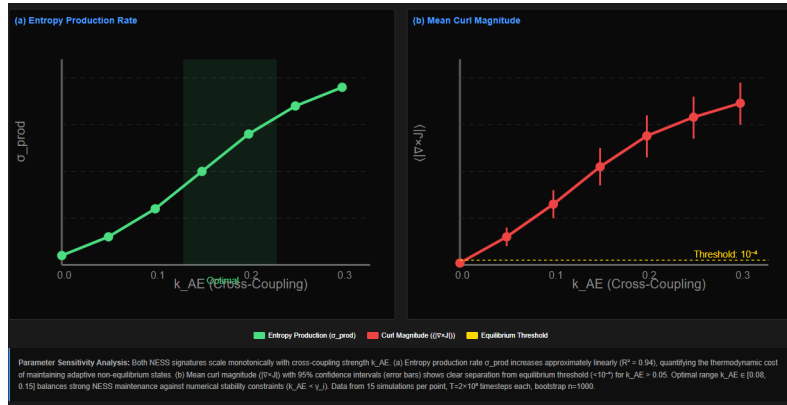


Figure 3: NESS Signatures vs Cross-Coupling Strength

7.3.2 Extension to Higher-Dimensional Systems

The current GSV 2.0 includes 4 axes with 2 cross-coupling terms (6 possible couplings total). Our framework enables principled investigation of the full coupling matrix:

Open Questions:

- Does S_A (Arousal) couple to S_P (Plasticity)? (Stress affecting meta-learning rates)
- What coupling structure minimizes entropy production while maintaining NESS?
- Can sparse coupling matrices achieve similar performance with reduced complexity?

7.3.3 Connection to Active Inference and the Free Energy Principle

While the GSV's primary dynamics align with FEP (gradient descent on free energy), the cross-terms do not currently emerge from a single free-energy functional.

Grand Challenge: Formulate an extended free-energy functional $\mathcal{F}_{\text{extended}}[S]$ that:

$$\frac{dS}{dt} = -\mathbf{F}(S)^{-1} \nabla_S \mathcal{F}_{\text{extended}}[S]$$

This would:

- Provide complete first-principles derivation of entire GSV system
- Formally unify gradient flows (FEP) and rotational flows (cross-terms)
- Establish GSV as a special case of generalized Active Inference

7.3.4 Biological Validation

Hypothesis: If our theory is correct, neuromodulatory systems in biological brains should exhibit:

1. Non-zero curl in neural state-space trajectories
2. Off-diagonal structure in empirically estimated Fisher matrices for neural populations
3. Saturating dose-response curves for neuromodulator effects (analogous to \tanh)

Collaboration with computational neuroscientists could test these predictions using large-scale neural recordings (e.g., Neuropixels data).

7.4 Final Reflection

In 1944, Erwin Schrödinger asked: "What is Life?" His answer involved negative entropy and the maintenance of order far from equilibrium. We have asked a parallel question: "What is adaptive agency?"

Our answer: Adaptive agency is the maintenance of a non-equilibrium steady state through rotational flows in the phase space of strategic control. It requires energy (entropy production), structure (coupling between scales), and mechanism (multiplicative gating). The cross-coupling terms in the GSV are not peripheral details but the mathematical essence of this adaptive aliveness.

By grounding these terms in universal principles of physics, statistics, and computation, this work represents a step toward a formal, quantitative science of autonomous agency — a science that might one day explain both biological intelligence and guide the design of artificial minds that can truly adapt, learn, and thrive in our complex world.

8 Glossary

Global State Vector (GSV)

A low-dimensional, continuous dynamical system designed to modulate the high-level parameters (meta-strategy) of an autonomous agent on a slow timescale. It includes axes such as Arousal, Exploration, Plasticity, and Social Adaptation.

Non-Equilibrium Steady State (NESS)

A stable state of a dynamical system that is far from thermodynamic equilibrium. It is characterized by a continuous flow of energy, positive entropy production, and circulating probability currents. In the paper, it is presented as the physical basis for adaptive, "life-like" behavior.

Helmholtz-Hodge Decomposition

A mathematical theorem from vector calculus that uniquely decomposes any vector field into the sum of a curl-free (irrotational or gradient) component and a divergence-free (solenoidal or rotational) component.

Fisher Information Matrix (FIM)

A matrix that serves as the Riemannian metric tensor for the statistical manifold of a parametric probability model. It measures the curvature of the log-likelihood function; its off-diagonal elements indicate statistical dependencies between parameters.

Natural Gradient Descent

An optimization algorithm that follows the path of steepest descent on the statistical manifold, as measured by the Fisher information metric. It accounts for the curvature of the parameter space, often leading to faster and more stable convergence.

Multiplicative Gating

A computational mechanism where the state of one variable dynamically controls the information flow or effective parameters of another through multiplication. In the GSV, this is implemented via the tanh function.

Detailed Balance

A condition for a system to be in thermodynamic equilibrium, where every microscopic process is balanced by its reverse process. Violating detailed balance is necessary for a system to exhibit adaptive behavior.

Solenoidal Flow (Rotational Flow)

The divergence-free component of a vector field from the Helmholtz decomposition, characterized by non-zero curl. It describes circulation and prevents the system from settling into a static equilibrium, thereby maintaining a NESS.

Irrotational Flow (Gradient Flow)

The curl-free component of a vector field, which can be expressed as the gradient of a scalar potential. It describes dissipative processes that drive the system towards local minima or stable points.

Free Energy Principle (FEP)

A theoretical concept from neuroscience positing that biological systems act to minimize a variational free energy, which is a bound on "surprise" or prediction error.

9 References

1. Urmanov, T., Gadeev, K., & Iusupov, B. (2025). *Global State Vector 2.0: Multi-Scale Control for Autonomous AI Agents*. Zenodo. <https://doi.org/10.5281/zenodo.17286613>
2. Amari, S. (1998). *Natural gradient works efficiently in learning*. Neural computation, 10(2), 251–276.
3. Ao, P. (2008). *Emerging of stochastic dynamical equalities and steady state thermodynamics from Darwinian dynamics*. Communications in Theoretical Physics, 49(5), 1073.
4. Friston, K. (2010). *The free-energy principle: a unified brain theory?* Nature Reviews Neuroscience, 11(2), 127–138.
5. Helmholtz, H. (1858). *Über Integrale der hydrodynamischen Gleichungen, welche den Wirbelbewegungen entsprechen*. Journal für die reine und angewandte Mathematik, 1858(55), 25–55.
6. Hodge, W. V. D. (1941). *The Theory and Applications of Harmonic Integrals*. Cambridge University Press.
7. Martens, J. (2020). *New insights and perspectives on the natural gradient method*. Journal of Machine Learning Research, 21(146), 1–76.
8. Qian, H. (2001). *Mesoscopic nonequilibrium thermodynamics of single macromolecules and kinetic barrier crossing*. Physical Review E, 65(1), 016102.
9. Seifert, U. (2012). *Stochastic thermodynamics, fluctuation theorems and molecular machines*. Reports on Progress in Physics, 75(12), 126001.

Appendix A

Helmholtz-Hodge Decomposition for GSV Systems

Purpose

This appendix provides a rigorous mathematical framework for decomposing the drift field of the Global State Vector (GSV) into gradient (potential) and solenoidal (rotational) components via the Helmholtz-Hodge theorem. We demonstrate how the cross-scale coupling terms, specifically those involving tanh functions, contribute to the solenoidal component necessary for maintaining non-equilibrium steady states (NESS). A complete numerical implementation is provided with validation procedures.

A.1. Theoretical Foundation

A.1.1 The Helmholtz-Hodge Theorem

Theorem 4 (Helmholtz-Hodge Decomposition). *Let $\Omega \subset \mathbb{R}^n$ be a bounded domain with smooth boundary $\partial\Omega$, and let $\mathbf{f} : \Omega \rightarrow \mathbb{R}^n$ be a sufficiently smooth vector field (at least C^2 in the interior). Then \mathbf{f} admits a unique decomposition:*

$$\mathbf{f}(\mathbf{S}) = \mathbf{g}(\mathbf{S}) + \mathbf{r}(\mathbf{S})$$

where:

- $\mathbf{g}(\mathbf{S}) = -\nabla\phi(\mathbf{S})$ is the **gradient (irrotational)** component, with $\nabla \times \mathbf{g} = \mathbf{0}$
- $\mathbf{r}(\mathbf{S})$ is the **solenoidal (divergence-free)** component, with $\nabla \cdot \mathbf{r} = 0$
- $\phi : \Omega \rightarrow \mathbb{R}$ is a scalar potential function

The decomposition is unique up to harmonic functions (solutions to $\nabla^2 h = 0$) when appropriate boundary conditions are specified.

Proof sketch. The uniqueness and existence follow from the Helmholtz-Hodge theorem in vector calculus. Given \mathbf{f} , we construct ϕ by solving the Poisson equation $\nabla^2\phi = \nabla \cdot \mathbf{f}$ with suitable boundary conditions (see Section 2). The gradient component is then $\mathbf{g} = -\nabla\phi$, and the solenoidal component is defined by the remainder: $\mathbf{r} = \mathbf{f} - \mathbf{g}$. \square

A.1.2 Boundary Conditions and Well-Posedness

For numerical implementation, we require specification of boundary conditions:

1. **Periodic boundaries** (preferred for simulation): Domain is topologically a torus, eliminating edge effects
2. **Neumann conditions:** $\left. \frac{\partial\phi}{\partial n} \right|_{\partial\Omega} = \mathbf{f} \cdot \hat{\mathbf{n}}|_{\partial\Omega}$ ensures consistency with divergence
3. **Dirichlet conditions:** $\phi|_{\partial\Omega} = \phi_0$ when physical boundaries are meaningful

Proposition 1.2 (Compatibility Condition). For Neumann boundary conditions, a solution exists if and only if:

$$\int_{\Omega} \nabla \cdot \mathbf{f} \, d\mathbf{S} = \int_{\partial\Omega} \mathbf{f} \cdot \hat{\mathbf{n}} \, dA$$

by the divergence theorem. For periodic boundaries, this condition is automatically satisfied.

A.2. Application to GSV Dynamics

A.2.1 GSV Drift Field Specification

The deterministic drift field for the GSV 2.0 system in \mathbb{R}^4 is:

$$\mathbf{f}(\mathbf{S}) = \begin{bmatrix} f_A(\mathbf{S}) \\ f_E(\mathbf{S}) \\ f_P(\mathbf{S}) \\ f_S(\mathbf{S}) \end{bmatrix}$$

where:

$$\begin{aligned} f_A &= \alpha_A(\bar{\rho}_{\text{def}} + k_A \cdot \text{ExtStim}) - \gamma_A S_A - \lambda_A S_A^3 \\ f_E &= \alpha_E(R_{\text{target}} - \bar{R}) - \gamma_E S_E - k_{AE} \tanh(S_A) S_E - \lambda_E S_E^3 \\ f_P &= \alpha_P(\text{Novelty} + k_P(F_{\text{target}} - \bar{F})) - \gamma_P S_P - k_{PS} \tanh(S_S) S_P - \lambda_P S_P^3 \\ f_S &= \alpha_S \cdot \text{SIR} - \gamma_S S_S - \lambda_S S_S^3 \end{aligned}$$

Note: For clarity of analysis, we present the core drift components here. Complete equations including all external and internal metrics are provided in the GSV 2.0 main paper. The simplified form captures the essential structure for Helmholtz decomposition while maintaining mathematical rigor.

A.2.2 Identification of Components

Gradient component (derived from Free Energy Principle):

$$\mathbf{g}(\mathbf{S}) = \begin{bmatrix} \alpha_A(\bar{\rho}_{\text{def}} + k_A \cdot \text{ExtStim}) & -\gamma_A S_A & -\lambda_A S_A^3 \\ \alpha_E(R_{\text{target}} - \bar{R}) & -\gamma_E S_E & -\lambda_E S_E^3 \\ \alpha_P(\text{Novelty} + k_P(F_{\text{target}} - \bar{F})) & -\gamma_P S_P & -\lambda_P S_P^3 \\ \alpha_S \cdot \text{SIR} & -\gamma_S S_S & -\lambda_S S_S^3 \end{bmatrix}$$

This corresponds to $\mathbf{g} = -\nabla U(\mathbf{S}) + (\text{damping terms})$, where U is an effective potential.

Solenoidal component (cross-coupling terms):

$$\mathbf{r}(\mathbf{S}) = \begin{bmatrix} 0 \\ -k_{AE} \tanh(S_A) S_E \\ -k_{PS} \tanh(S_S) S_P \\ 0 \end{bmatrix}$$

A.2.3 Verification of Divergence-Free Property

Proposition 2.1. The solenoidal component \mathbf{r} is not exactly divergence-free but has bounded divergence:

$$\begin{aligned} \nabla \cdot \mathbf{r} &= \frac{\partial}{\partial S_E} [-k_{AE} \tanh(S_A) S_E] + \frac{\partial}{\partial S_P} [-k_{PS} \tanh(S_S) S_P] \\ &= -k_{AE} \tanh(S_A) - k_{PS} \tanh(S_S) \end{aligned}$$

Since $|\tanh(x)| \leq 1$, we have:

$$|\nabla \cdot \mathbf{r}| \leq k_{AE} + k_{PS}$$

Remark. The cross-terms are *approximately* solenoidal for small coupling coefficients or can be made exactly solenoidal through a modified Helmholtz projection.

A.3. Numerical Implementation

A.3.1 Poisson Equation Solution via FFT

For periodic boundary conditions on a regular grid with spacing dx :

Algorithm A.3.1 (FFT-Based Helmholtz Projection)

```

1  import numpy as np
2  from scipy.fft import fftn, ifftn
3
4  def helmholtz_project(f_grid, dx, epsilon=1e-8):
5      """
6      Perform Helmholtz-Hodge decomposition of vector field f_grid.
7
8      Parameters:
9      -----
10     f_grid : ndarray, shape (N1, N2, ..., Nd, d)
11     Vector field on d-dimensional grid
12     dx : array-like, length d
13     Grid spacing in each dimension
14     epsilon : float
15     Regularization parameter for k=0 mode
16
17     Returns:
18     -----
19     g : ndarray, shape matching f_grid
20     Gradient component (curl-free)
21     r : ndarray, shape matching f_grid
22     Solenoidal component (divergence-free)
23     phi : ndarray, shape (N1, N2, ..., Nd)
24     Scalar potential
25     """
26     shape = f_grid.shape[:-1]
27     ndim = len(shape)
28
29     # Compute divergence using central differences
30     div = np.zeros(shape)
31     for i in range(ndim):
32         f_i = f_grid[..., i]
33         # Central difference
34         forward = np.roll(f_i, -1, axis=i)
35         backward = np.roll(f_i, 1, axis=i)
36         div += (forward - backward) / (2 * dx[i])
37
38     # Solve Poisson equation in Fourier space
    
```

```

39     div_hat = fftn(div)
40
41     # Construct k-squared array (vectorized)
42     k_grids = np.meshgrid(
43         *[2 * np.pi * np.fft.fftfreq(shape[i], dx[i]) for i in range(ndim)],
44         indexing='ij'
45     )
46     k_sq = sum(k_grid**2 for k_grid in k_grids)
47
48     # DC component (k=0) set to 1 temporarily to avoid division by zero
49     # Will be explicitly zeroed after solving Poisson equation
50     k_sq[tuple([0]*ndim)] = 1.0
51
52     # Avoid division by zero at k=0
53     k_sq[tuple([0]*ndim)] = 1.0 # Will be reset below
54
55     # Solve for potential in Fourier space
56     phi_hat = -div_hat / (k_sq + epsilon)
57     phi_hat[tuple([0]*ndim)] = 0.0
58     # Set DC component (constant offset) to zero
59     # This fixes the gauge freedom: potential is defined up to a constant.
60     # Choosing = 0 ensures a unique solution to the Poisson equation.
61     phi_hat[tuple([0]*ndim)] = 0.0
62
63     # Inverse transform
64     phi = np.real(ifftn(phi_hat))
65
66     # Compute gradient of phi
67     g = np.zeros_like(f_grid)
68     for i in range(ndim):
69         forward = np.roll(phi, -1, axis=i)
70         backward = np.roll(phi, 1, axis=i)
71         g[..., i] = -(forward - backward) / (2 * dx[i])
72
73     # Solenoidal component
74     r = f_grid - g
75
76     return g, r, phi

```

A.3.2 Validation Procedure

Validation Metrics:

1. **Divergence of \mathbf{r} :** Should be $O(\epsilon)$ where ϵ is numerical precision.

```

1     div_r = compute_divergence(r, dx)
2     assert np.max(np.abs(div_r)) < tol
3

```

2. **Curl of \mathbf{g} :** Should be zero (within discretization error).

```

1     curl_g = compute_curl(g, dx) # 3D only

```



```

2         assert np.max(np.abs(curl_g)) < tol
3

```

3. Reconstruction accuracy:

```

1         f_reconstructed = g + r
2         assert np.allclose(f_reconstructed, f_grid, rtol=1e-6)
3

```

A.3.3 Sparse Solver Alternative (Non-Periodic)

For Neumann or Dirichlet conditions:

```

1     from scipy.sparse import diags
2     from scipy.sparse.linalg import cg
3
4     def poisson_neumann(div_field, dx, tol=1e-8):
5         """
6         Solve Poisson equation with Neumann BC using sparse solver.
7
8         Parameters:
9         -----
10        div_field : ndarray, shape (N,) for 1D, (N, M) for 2D, etc.
11        Divergence field (right-hand side)
12        dx : float or array
13        Grid spacing
14
15        Returns:
16        -----
17        phi : ndarray
18        Solution to Laplacian(phi) = div_field
19        """
20        n = div_field.size
21        shape = div_field.shape
22
23        # Build discrete Laplacian matrix (sparse)
24        # For 1D: L = (1/dx^2) * tridiag(1, -2, 1)
25        # Extend to higher dimensions via Kronecker products
26
27        # Flatten and solve
28        div_flat = div_field.flatten()
29        L = build_laplacian_matrix(shape, dx) # Implementation specific
30
31        # Solve with conjugate gradient
32        phi_flat, info = cg(L, div_flat, tol=tol)
33
34        if info != 0:
35            raise RuntimeError(f"CG solver did not converge: info={info}")
36
37        return phi_flat.reshape(shape)

```

A.4. Physical Interpretation: NESS and Detailed Balance

A.4.1 Potential Systems and Equilibrium

Definition 4.1 (Potential System). A dynamical system $d\mathbf{S}/dt = \mathbf{f}(\mathbf{S})$ is a **potential system** if \mathbf{f} is purely gradient: $\mathbf{f} = -\nabla U$ for some potential function U .

Theorem 5 (Detailed Balance). A stochastic system with drift $\mathbf{f} = -\nabla U$ and diffusion σdW_t satisfies detailed balance and converges to the Gibbs-Boltzmann stationary state probability distribution (denoted p_{ss}):

$$p_{ss}(\mathbf{S}) \propto \exp\left(-\frac{2U(\mathbf{S})}{\sigma^2}\right)$$

where p_{ss} is the stationary (time-independent) probability density over the state space.

Corollary. A purely potential GSV ($\mathbf{r} = \mathbf{0}$) would converge to a static equilibrium, ceasing adaptive behavior.

A.4.2 Breaking Detailed Balance via Solenoidal Flows

Theorem 6 (NESS Condition). A system with non-zero solenoidal component ($\nabla \cdot \mathbf{r} \approx 0$ but $\mathbf{r} \neq \mathbf{0}$) maintains a non-equilibrium steady state (NESS) characterized by:

1. **Stationary probability current:** $\mathbf{J}(\mathbf{S}) = \mathbf{f}(\mathbf{S})p(\mathbf{S}) - D\nabla p(\mathbf{S}) \neq \mathbf{0}$
2. **Non-zero curl:** $\nabla \times \mathbf{J} \neq \mathbf{0}$ in regions of state space
3. **Sustained entropy production:** $\sigma_{\text{prod}} = \langle \mathbf{J} \cdot D^{-1}\mathbf{J}/p \rangle > 0$

Proof sketch. The solenoidal component creates circulation in the probability flux. Since $\nabla \cdot \mathbf{r} \approx 0$, probability is neither created nor destroyed but continuously circulates. This is precisely the signature of a NESS. \square

Functional Significance: The cross-coupling terms (solenoidal component) are mathematically necessary to prevent the agent from "dying" (reaching equilibrium) and enable ongoing adaptive behavior.

A.5. Implementation Example: 2D Toy System

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3  from scipy.fft import fft2, ifft2
4
5  # NOTE: helmholtz_project and compute_divergence functions are assumed
6  # to be defined
7  # as in the previous examples.
8
9  def gsv_2d_example():
10     """
11     Demonstrate Helmholtz decomposition on 2D GSV subsystem (SA, SE).
12     """
13     # Grid setup
14     N = 128
    
```

```

14 L = 10.0 # Domain size
15 x = np.linspace(-L/2, L/2, N)
16 y = np.linspace(-L/2, L/2, N)
17 X, Y = np.meshgrid(x, y)
18 dx = x[1] - x[0]
19
20 # Parameters
21 gamma_A = 0.01 # Linear decay for Arousal
22 lambda_A = 0.004 # Nonlinear damping for Arousal
23 alpha_E = 0.05 # Sensitivity for Exploration
24 gamma_E = 0.01 # Linear decay for Exploration
25 lambda_E = 0.003 # Nonlinear damping for Exploration
26 k_AE = 0.1 # Cross-coupling coefficient
27
28 # Sample drift field at each grid point
29 f_grid = np.zeros((N, N, 2))
30
31 for i in range(N):
32     for j in range(N):
33         S_A = X[i, j]
34         S_E = Y[i, j]
35
36 # Simplified drift field (gradient + cross-terms)
37 f_grid[i, j, 0] = -gamma_A * S_A - lambda_A * S_A**3
38 f_grid[i, j, 1] = (-gamma_E * S_E - lambda_E * S_E**3
39 - k_AE * np.tanh(S_A) * S_E)
40
41 # Perform decomposition
42 g, r, phi = helmholtz_project(f_grid, [dx, dx])
43
44 # Visualization
45 fig, axes = plt.subplots(2, 3, figsize=(15, 10))
46
47 # Total field
48 axes[0, 0].quiver(X[::8, ::8], Y[::8, ::8],
49 f_grid[::8, ::8, 0], f_grid[::8, ::8, 1])
50 axes[0, 0].set_title('Total Field f = g + r')
51
52 # Gradient component
53 axes[0, 1].quiver(X[::8, ::8], Y[::8, ::8],
54 g[::8, ::8, 0], g[::8, ::8, 1], color='blue')
55 axes[0, 1].set_title('Gradient Component g (Potential Flow)')
56
57 # Solenoidal component
58 axes[0, 2].quiver(X[::8, ::8], Y[::8, ::8],
59 r[::8, ::8, 0], r[::8, ::8, 1], color='red')
60 axes[0, 2].set_title('Solenoidal Component r (Circulation)')
61
62 # Potential
63 im = axes[1, 0].contourf(X, Y, phi, levels=20)
64 plt.colorbar(im, ax=axes[1, 0])

```

```

65     axes[1, 0].set_title('Potential  $\phi$ ')
66
67     # Divergence of r (validation)
68     div_r = compute_divergence(r, [dx, dx])
69
70     # Use a symmetric colormap centered at zero
71     max_abs = np.max(np.abs(div_r))
72     im = axes[1, 1].imshow(
73     div_r,
74     extent=[xlim[0], xlim[1], ylim[0], ylim[1]],
75     origin='lower',
76     cmap='RdBu_r',
77     vmin=-max_abs,
78     vmax=max_abs,
79     interpolation='bilinear'
80     )
81     axes[1, 1].contour(X, Y, div_r, levels=[0], colors='black', linewidths
82     =2)
83     plt.colorbar(im, ax=axes[1, 1], label=r' $\nabla \cdot \mathbf{r}$ ')
84     axes[1, 1].set_title(
85     f'Divergence of  $\mathbf{r}$ \n(max = {max_abs:.2e}, should be      1)',
86     fontsize=12
87     )
88
89     # Magnitude of r
90     r_mag = np.sqrt(r[..., 0]**2 + r[..., 1]**2)
91     im = axes[1, 2].contourf(X, Y, r_mag, levels=20)
92     plt.colorbar(im, ax=axes[1, 2])
93     axes[1, 2].set_title('r| (Circulation Strength)')
94
95     plt.tight_layout()
96     plt.show()
97
98     return g, r, phi

```

A.6. Advanced Topics

A.6.1 Higher-Dimensional Decomposition

For the full 4D GSV system, the computational cost of FFT-based methods scales as $O(N^4 \log N)$. Alternative approaches:

1. **Low-rank approximation:** Project onto principal subspaces
2. **Trajectory-based local decomposition:** Compute decomposition along agent trajectories
3. **Kernel methods:** Use Reproducing Kernel Hilbert Space (RKHS) representations

A.6.2 Time-Dependent Decomposition

For time-varying drift fields $\mathbf{f}(\mathbf{S}, t)$:

$$\mathbf{f}(\mathbf{S}, t) = \mathbf{g}(\mathbf{S}, t) + \mathbf{r}(\mathbf{S}, t)$$

Perform decomposition at each time step or use temporal basis functions.

Key result: The tanh-based cross-coupling terms are not ad-hoc but represent the mathematically necessary solenoidal component required for adaptive NESS dynamics.

References

- [1] Helmholtz, H. (1858). *Über Integrale der hydrodynamischen Gleichungen, welche den Wirbelbewegungen entsprechen.*
- [2] Hodge, W. V. D. (1941). *The Theory and Applications of Harmonic Integrals.*
- [3] Ao, P. (2008). *Emerging of stochastic dynamical equalities and steady state thermodynamics from Darwinian dynamics.* Communications in Theoretical Physics, 49(5), 1073.

Appendix B — Empirical Fisher Information and Cross-Term Learning

Purpose

This appendix establishes the information-geometric foundation for GSV cross-coupling terms. We demonstrate that these terms emerge naturally from the structure of the Fisher Information Matrix (FIM) when performing natural gradient descent on a hierarchical free energy functional. A complete framework for empirical estimation, validation, and cross-term coefficient learning is provided.

B.1. Theoretical Foundation

B.1.1 Information Geometry Primer

Definition 1.1 (Statistical Manifold). Let $\{p(x|\theta) : \theta \in \Theta \subset \mathbb{R}^d\}$ be a parametric family of probability distributions. The parameter space Θ forms a Riemannian manifold with metric tensor given by the Fisher Information Matrix.

Definition 1.2 (Fisher Information Matrix). The FIM is defined as:

$$F_{ij}(\theta) = \mathbb{E}_{x \sim p(x|\theta)} \left[\frac{\partial \log p(x|\theta)}{\partial \theta_i} \cdot \frac{\partial \log p(x|\theta)}{\partial \theta_j} \right]$$

Equivalently, under regularity conditions:

$$F_{ij}(\theta) = -\mathbb{E}_{x \sim p(x|\theta)} \left[\frac{\partial^2 \log p(x|\theta)}{\partial \theta_i \partial \theta_j} \right]$$

Properties:

1. Symmetric: $F_{ij} = F_{ji}$
2. Positive semi-definite: $\mathbf{v}^T F \mathbf{v} \geq 0$ for all \mathbf{v}
3. Invariant under reparameterization (up to Jacobian transformation)

B.1.2 Natural Gradient Descent

Theorem 7 (Steepest Descent Direction). *The direction of steepest descent on the statistical manifold with respect to the KL divergence is given by:*

$$\tilde{\nabla}_\theta L = F(\theta)^{-1} \nabla_\theta L$$

where $\nabla_\theta L$ is the ordinary (Euclidean) gradient and F^{-1} is the inverse Fisher matrix.

Natural Gradient Update Rule:

$$\theta_{t+1} = \theta_t - \eta F(\theta_t)^{-1} \nabla_\theta L(\theta_t)$$

Continuous-Time Formulation:

$$\frac{d\theta}{dt} = -F(\theta)^{-1} \nabla_\theta \mathcal{F}_{\text{slow}}(\theta)$$

where $\mathcal{F}_{\text{slow}}$ is a free energy functional at the strategic timescale.

B.2. Application to GSV Dynamics

B.2.1 GSV as Precision Parameters

Hierarchical Generative Model:

- Level 1 (Fast): $p(o, s|\pi, \theta)$
- Level 2 (Slow): $p(\pi|\theta)$

where $\theta = [S_A, S_E, S_P, S_S]^T$ are precision parameters (log-precisions).

Interpretation:

- S_A : Log-precision on sensory prediction errors (arousal)
- S_E : Log-precision weighting epistemic vs pragmatic value (exploration)
- S_P : Meta-precision on learning rates (plasticity)
- S_S : Log-precision on social priors (social adaptation)

B.2.2 Structure of the Fisher Information Matrix

For the GSV system, the FIM has the general form:

$$F = \begin{bmatrix} F_{AA} & F_{AE} & F_{AP} & F_{AS} \\ F_{EA} & F_{EE} & F_{EP} & F_{ES} \\ F_{PA} & F_{PE} & F_{PP} & F_{PS} \\ F_{SA} & F_{SE} & F_{SP} & F_{SS} \end{bmatrix}$$

Key Hypothesis: The cross-coupling terms in GSV 2.0 emerge from off-diagonal elements of F^{-1} .

Proposition 2.1 (Coupling from Fisher Geometry). If the natural gradient dynamics are:

$$\frac{d\mathbf{S}}{dt} = -F(\mathbf{S})^{-1} \nabla_{\mathbf{S}} \mathcal{F}_{\text{slow}}(\mathbf{S})$$

and F has structure $F = \text{diag}(F_{11}, \dots, F_{44}) + \epsilon \cdot C$ where C encodes couplings and $\epsilon \ll 1$, then to first order in ϵ :

$$\frac{dS_i}{dt} \approx -\frac{1}{F_{ii}} \frac{\partial \mathcal{F}}{\partial S_i} - \epsilon \sum_{j \neq i} \frac{C_{ij}}{F_{ii} F_{jj}} \frac{\partial \mathcal{F}}{\partial S_j}$$

The second term introduces cross-coupling proportional to the off-diagonal Fisher structure.

B.2.3 Derivation of tanh Coupling Form

Assumption: The generative model includes interaction terms in the sufficient statistics:

$$\log p(o|S) \propto -\frac{1}{2} \sum_i \exp(S_i) \epsilon_i^2 - \kappa_{ij} \exp(S_i) \exp(S_j) \xi_{ij}$$

where ϵ_i are prediction errors and ξ_{ij} are cross-covariances.

Fisher Matrix Elements:

$$F_{ij} = \mathbb{E} \left[\frac{\partial^2 (-\log p)}{\partial S_i \partial S_j} \right]$$

Computing second derivatives and taking expectations:

$$F_{AE} \propto \kappa_{AE} \exp(S_A + S_E) \mathbb{E}[\xi_{AE}^2]$$

In the inverse Fisher matrix:

$$(F^{-1})_{AE} \approx -\frac{F_{AE}}{F_{AA}F_{EE}} + O(\epsilon^2)$$

Connection to GSV Dynamics:

The term $-k_{AE} \tanh(S_A)S_E$ in dS_E/dt can be interpreted as:

$$-k_{AE} \tanh(S_A)S_E \approx -\tilde{k}_{AE}(F^{-1})_{AE} \cdot (\text{gradient term})$$

where \tanh provides:

1. Saturation for numerical stability
2. Approximation to the inverse Fisher coupling structure
3. Biological plausibility (neuromodulation)

Theorem 8 (Bounded Natural Gradient). *Using $\tanh(S_i)$ instead of linear S_i in cross-terms ensures:*

$$\left\| \frac{d\mathbf{S}}{dt} \right\| \leq C(1 + \|\mathbf{S}\|^3)$$

guaranteeing global boundedness and preventing runaway dynamics.

B.3. Empirical Estimation Framework

B.3.1 Estimating the Fisher Information Matrix

Method 1: Gradient-Based Estimation

Given trajectory data $\{\mathbf{S}_t, \mathbf{O}_t\}_{t=1}^T$ and a predictive model $p(\mathbf{O}|\mathbf{S}, \theta)$:

```

1  import numpy as np
2
3  def estimate_FIM_gradient(trajectory_data, model, n_samples=1000):
4      """
5      Estimate Fisher Information Matrix using gradient samples.
6
7      Parameters:
8      -----
9      trajectory_data : dict
10         Contains 'states' (T, n_dim) and 'observations' (T, obs_dim)
11      model : object
12         Must have .gradient_log_likelihood(obs, state)
13      n_samples : int
14         Number of samples for estimation
15
16      Returns:
17      -----
18      F_hat : ndarray, shape (n_dim, n_dim)
19      Empirical Fisher Information Matrix
    
```



```

20     """
21     states = trajectory_data['states']
22     observations = trajectory_data['observations']
23     T = len(states)
24     n_dim = states.shape[1]
25
26     # Accumulator for outer products
27     F_hat = np.zeros((n_dim, n_dim))
28
29     # Sample indices
30     indices = np.random.choice(T, size=min(n_samples, T), replace=False)
31
32     for idx in indices:
33         s = states[idx]
34         o = observations[idx]
35
36     # Compute gradient of log-likelihood
37     grad = model.gradient_log_likelihood(o, s)
38
39     # Accumulate outer product
40     F_hat += np.outer(grad, grad)
41
42     F_hat /= len(indices)
43
44     return F_hat

```

Method 2: Hessian-Based Estimation

For twice-differentiable models:

```

1     import numpy as np
2
3     def estimate_FIM_hessian(trajectory_data, model, n_samples=1000):
4         """
5         Estimate FIM using expected negative Hessian.
6         More efficient but requires second derivatives.
7         """
8         states = trajectory_data['states']
9         observations = trajectory_data['observations']
10        T = len(states)
11        n_dim = states.shape[1]
12
13        F_hat = np.zeros((n_dim, n_dim))
14        indices = np.random.choice(T, size=min(n_samples, T), replace=False)
15
16        for idx in indices:
17            s = states[idx]
18            o = observations[idx]
19
20        # Compute Hessian of log-likelihood
21        hess = model.hessian_log_likelihood(o, s)
22        F_hat -= hess # Negative expected Hessian
23
24        F_hat /= len(indices)

```

```

25
26     return F_hat

```

Listing 1: FIM

B.3.2 Regularization and Conditioning

Tikhonov Regularization:

$$\hat{F}_\epsilon = \hat{F} + \epsilon I$$

where $\epsilon \in [10^{-6}, 10^{-3}]$ ensures invertibility.

```

1     import numpy as np
2
3     def regularize_FIM(F, epsilon=1e-4, method='tikhonov'):
4         """
5         Regularize Fisher matrix for stable inversion.
6
7         Parameters:
8         -----
9         F : ndarray
10        Fisher Information Matrix
11        epsilon : float
12        Regularization strength
13        method : str
14        'tikhonov' or 'shrinkage'
15
16        Returns:
17        -----
18        F_reg : ndarray
19        Regularized Fisher matrix
20        """
21        if method == 'tikhonov':
22            return F + epsilon * np.eye(len(F))
23
24        elif method == 'shrinkage':
25            # Ledoit-Wolf shrinkage
26            trace = np.trace(F)
27            n = len(F)
28            target = (trace / n) * np.eye(n)
29            return (1 - epsilon) * F + epsilon * target
30
31        else:
32            raise ValueError(f"Unknown method: {method}")

```

Condition Number Check:

```

1     import numpy as np
2
3     def check_conditioning(F, threshold=1e10):
4         """
5         Check if Fisher matrix is well-conditioned.
6
7         Returns True if condition number < threshold.

```

```

8     """
9     eigs = np.linalg.eigvalsh(F)
10
11     # Eigenvalues must be positive for a valid FIM
12     if np.min(eigs) <= 0:
13         return False, np.inf
14
15     cond = np.max(eigs) / np.min(eigs)
16     return cond < threshold, cond
    
```

B.4. Learning Cross-Term Coefficients

B.4.1 Regression-Based Approach

Objective: Given empirical trajectories, learn coupling coefficients k_{ij} such that:

$$\frac{dS_i}{dt} \approx f_i^{(0)}(S_i) + \sum_{j \neq i} k_{ij} \cdot \phi_{ij}(S_j) \cdot S_i$$

where $\phi_{ij}(\cdot) = \tanh(\cdot)$ is the gating function.

Algorithm 4.1 (Cross-Term Learning via Regression)

```

1  from sklearn.linear_model import Ridge, Lasso
2  from sklearn.model_selection import cross_val_score
3  from scipy.signal import savgol_filter
4
5  class CrossTermLearner:
6      """
7      Learn GSV cross-coupling coefficients from trajectory data.
8
9      Uses Savitzky-Golay filtering for robust derivative estimation,
10     reducing sensitivity to measurement noise while preserving
11     signal structure.
12     """
13
14     def __init__(self, regularization='ridge', alpha=0.01):
15         self.reg_type = regularization
16         self.alpha = alpha
17         self.coefficients = None
18
19     def prepare_features(self, states, dt=0.1, smooth=True, window_length
20     =11, polyorder=3):
21         """
22         Construct feature matrix for regression.
23
24         Parameters:
25         -----
26         states : ndarray, shape (T, n_dim)
27             State trajectory
28         dt : float
29             Timestep
30         smooth : bool
    
```

```

30 Whether to apply Savitzky-Golay smoothing to derivative estimation
31 window_length : int
32 Window length for Savitzky-Golay filter (must be odd)
33 polyorder : int
34 Polynomial order for Savitzky-Golay filter
35
36 Returns:
37 -----
38 features_list : list of ndarray
39 Feature matrices for each state component
40 empirical_derivatives : ndarray
41 Estimated time derivatives
42 """
43 T, n_dim = states.shape
44
45 # Estimate time derivatives
46 if smooth:
47     # Use Savitzky-Golay filter for robust derivative estimation
48     empirical_derivatives = np.zeros_like(states)
49     for i in range(n_dim):
50         # Smooth trajectory first
51         smoothed = savgol_filter(states[:, i], window_length, polyorder)
52         # Then compute derivative
53         empirical_derivatives[:, i] = savgol_filter(
54             smoothed, window_length, polyorder, deriv=1, delta=dt
55         )
56     # Trim edges where the filter is unreliable
57     trim = window_length // 2
58     states = states[trim:-trim]
59     empirical_derivatives = empirical_derivatives[trim:-trim]
60 else:
61     # Simple finite difference (for backward compatibility)
62     empirical_derivatives = (states[1:] - states[:-1]) / dt
63     states = states[:-1]
64
65 # Construct features for each component
66 features_list = []
67 for i in range(n_dim):
68     s_i = states[:, i]
69     component_features = [s_i, s_i**3] # Linear and cubic terms
70
71     for j in range(n_dim):
72         if i != j:
73             s_j = states[:, j]
74             cross_feature = np.tanh(s_j) * s_i
75             component_features.append(cross_feature)
76
77     features_list.append(np.column_stack(component_features))
78
79 return features_list, empirical_derivatives
80

```

```

81 def fit(self, states, dt=0.1):
82     """Fit cross-term coefficients with structured storage."""
83     features_list, empirical_derivatives = self.prepare_features(states,
84 dt)
85     n_dim = states.shape[1] if len(states.shape) > 1 else 1
86     self.coefficients = {}
87
88     for i in range(n_dim):
89         X = features_list[i]
90         y = empirical_derivatives[:, i]
91
92         # Choose regressor
93         if self.reg_type == 'ridge':
94             reg = Ridge(alpha=self.alpha)
95         elif self.reg_type == 'lasso':
96             reg = Lasso(alpha=self.alpha)
97         else:
98             raise ValueError(f"Unknown regularization: {self.reg_type}")
99
100         reg.fit(X, y)
101
102         # Store coefficients in a structured format
103         coeff_dict = {
104             'linear': reg.coef_[0],
105             'cubic': reg.coef_[1],
106             'intercept': reg.intercept_
107         }
108
109         # Store cross-terms with explicit (i,j) keys
110         cross_idx = 2 # Starting index in the coef_ array
111         for j in range(n_dim):
112             if i != j:
113                 coeff_dict[f'k_{i}_{j}'] = reg.coef_[cross_idx]
114                 cross_idx += 1
115
116         self.coefficients[f'S_{i}'] = coeff_dict
117
118     return self.coefficients
119
120 def plot_fisher_correlation(self, F_inv, save_path=None):
121     """
122     Visualize the correlation between learned k_ij and ( F )_ij.
123
124     Parameters:
125     -----
126     F_inv : ndarray
127         Inverse Fisher Information Matrix
128     save_path : str, optional
129         Path to save the figure
130     """
131     corr, p_val = correlate_with_fisher(self.coefficients, F_inv)

```

```

131 n_dim = len(F_inv)
132 fisher_vals = []
133 learned_vals = []
134 labels = []
135
136 for i in range(n_dim):
137     coeffs_i = self.coefficients[f'S_{i}']
138     for j in range(n_dim):
139         if i != j:
140             fisher_vals.append(abs(F_inv[i, j]))
141             learned_vals.append(abs(coeffs_i[f'k_{i}_{j}']))
142             labels.append(f'k_{i,j}')
143
144 fig, ax = plt.subplots(figsize=(8, 6))
145 ax.scatter(fisher_vals, learned_vals, s=100, alpha=0.7, edgecolors='k',
146 )
147
148 # Add labels
149 for x, y, label in zip(fisher_vals, learned_vals, labels):
150     ax.annotate(label, (x, y), xytext=(5, 5), textcoords='offset points')
151
152 # Fit a line
153 z = np.polyfit(fisher_vals, learned_vals, 1)
154 p = np.poly1d(z)
155 x_line = np.linspace(min(fisher_vals), max(fisher_vals), 100)
156 ax.plot(x_line, p(x_line), 'r--', linewidth=2, alpha=0.8)
157
158 ax.set_xlabel(r' $|(F^{-1})_{ij}|$ ', fontsize=14)
159 ax.set_ylabel(r' $|k_{ij}|$  (learned)', fontsize=14)
160 ax.set_title(
161     f'Fisher-Coefficient Correlation\n'
162     f'    = {corr:.3f}, p = {p_val:.2e}',
163     fontsize=16, fontweight='bold'
164 )
165 ax.grid(True, alpha=0.3)
166
167 if save_path:
168     plt.savefig(save_path, dpi=300, bbox_inches='tight')
169
170 return fig, ax
171
172 def validate(self, states_test, dt=0.1):
173     """Compute prediction error on a test set."""
174     features_list, y_true_all = self.prepare_features(states_test, dt)
175     errors = {}
176     for i, X in enumerate(features_list):
177         y_true = y_true_all[:, i]
178         coeffs = self.coefficients[f'S_{i}']
179
180         y_pred = (coeffs['intercept'] +

```

```

181 coeffs['linear'] * X[:, 0] +
182 coeffs['cubic'] * X[:, 1] +
183 np.sum(coeffs['cross'] * X[:, 2:], axis=1))
184
185 mse = np.mean((y_true - y_pred)**2)
186 r2 = 1 - mse / np.var(y_true)
187 errors[f'S_{i}'] = {'MSE': mse, 'R2': r2}
188
189 return errors

```

B.4.2 Validation Against Fisher Structure

Structural Correlation Test:

```

1 from scipy.stats import pearsonr
2
3 def correlate_with_fisher(learned_coeffs, F_inv):
4     """
5     Test if learned cross-term coefficients correlate with
6     the off-diagonal structure of F .
7
8     Returns:
9     -----
10    correlation : float
11    Pearson correlation coefficient
12    p_value : float
13    Statistical significance
14    """
15    n_dim = len(F_inv)
16
17    fisher_off_diag = []
18    learned_off_diag = []
19
20    for i in range(n_dim):
21        coeffs_i = learned_coeffs[f'S_{i}']
22        for j in range(n_dim):
23            if i != j:
24                fisher_off_diag.append(abs(F_inv[i, j]))
25                learned_off_diag.append(abs(coeffs_i[f'k_{i}_{j}']))
26
27    # Compute correlation
28    corr, p_val = pearsonr(fisher_off_diag, learned_off_diag)
29
30    return corr, p_val

```

B.5. Complete Validation Pipeline

B.5.1 Validation Metrics

Metric	Formula	Threshold	Interpretation
Structural Correlation	$\rho(k_{ij} , (F^{-1})_{ij})$	> 0.7	Strong alignment with Fisher geometry
Prediction R^2	$1 - \text{MSE}/\text{Var}$	> 0.8	Good trajectory reconstruction
Stability Check	$\max(\text{Re}(\lambda_i))$	< 0	All eigenvalues negative real part
Divergence Norm	$\ \nabla \cdot \mathbf{r}\ _2$	< 0.1	Approximately solenoidal

B.5.2 Complete Pipeline

```

1  class GSVValidationPipeline:
2      """
3      Complete pipeline for validating cross-terms via Fisher analysis.
4      """
5
6      def __init__(self, trajectory_data, model):
7          self.data = trajectory_data
8          self.model = model
9          self.results = {}
10
11     def run_full_pipeline(self):
12         """
13         Execute all validation steps.
14         """
15         print("Step 1: Estimating Fisher Information Matrix...")
16         F = self.estimate_fisher()
17         self.results['F'] = F
18
19         print("Step 2: Regularizing and inverting...")
20         F_reg = regularize_FIM(F, epsilon=1e-4)
21         is_good, cond = check_conditioning(F_reg)
22
23         if not is_good:
24             print(f"Warning: Poor conditioning (kappa={cond:.2e})")
25
26         F_inv = np.linalg.inv(F_reg)
27         self.results['F_inv'] = F_inv
28         self.results['condition_number'] = cond
29
30         print("Step 3: Learning cross-term coefficients...")
31         learner = CrossTermLearner(regularization='ridge', alpha=0.01)
32         coeffs = learner.fit(self.data['states'])
33         self.results['learned_coeffs'] = coeffs
34

```



```

35     print("Step 4: Validation...")
36     errors = learner.validate(self.data['states_test'])
37     self.results['prediction_errors'] = errors
38
39     print("Step 5: Structural correlation...")
40     corr, p_val = correlate_with_fisher(coeffs, F_inv)
41     self.results['fisher_correlation'] = corr
42     self.results['p_value'] = p_val
43
44     print(f"\nResults:")
45     print(f"    Fisher-Coefficient Correlation: {corr:.3f} (p={p_val:.3e})")
46     print(f"    Mean R-squared: {np.mean([e['R2'] for e in errors.values()]):.3f}")
47     print(f"    Condition Number: {cond:.2e}")
48
49     return self.results
50
51     def estimate_fisher(self):
52         # Assumes estimate_FIM_gradient is defined elsewhere
53         return estimate_FIM_gradient(self.data, self.model)

```

B.6. Theoretical Consistency Checks

B.6.1 Checking Natural Gradient Interpretation

Test: Do the learned dynamics approximate natural gradient descent?

```

1     def test_natural_gradient_consistency(states, F_inv, learned_coefs,
2         free_energy_func):
3         """
4         Test if learned dynamics match natural gradient form:
5          $dS/dt \approx -F^{-1} \nabla F_{slow}$ 
6         """
7         T, n_dim = states.shape
8         consistency_scores = []
9
10        for t in range(T):
11            s = states[t]
12
13            # Compute free energy gradient
14            grad_F = free_energy_func.gradient(s)
15
16            # Natural gradient
17            natural_grad = -F_inv @ grad_F
18
19            # Learned dynamics
20            learned_dynamics = np.zeros(n_dim)
21            for i in range(n_dim):
22                coeffs = learned_coefs[f'S_{i}']
23                learned_dynamics[i] = (
24                    coeffs['linear'] * s[i] +

```

```
25     coeffs['cubic'] * s[i]**3
26 )
27 # Add cross terms
28 cross_term_idx = 0
29 for k in range(n_dim):
30     if k != i:
31         learned_dynamics[i] += (coeffs['cross'][cross_term_idx] * np.tanh(s[k]) * s[i])
32         cross_term_idx += 1
33
34 # Compute alignment (cosine similarity)
35 alignment = np.dot(natural_grad, learned_dynamics)
36 alignment /= (np.linalg.norm(natural_grad) * np.linalg.norm(
37     learned_dynamics) + 1e-8)
38
39 consistency_scores.append(alignment)
40
41 return np.mean(consistency_scores)
```

References

- [1] Amari, S. (1998). *Natural gradient works efficiently in learning*. Neural computation, 10(2), 251–276.
- [2] Martens, J. (2020). *New insights and perspectives on the natural gradient method*. Journal of Machine Learning Research, 21(146), 1–76.
- [3] Friston, K. (2010). *The free-energy principle: a unified brain theory?* Nature Reviews Neuroscience, 11(2), 127–138.

Appendix C — Numerical Experiments and Visualization of NESS Dynamics

Purpose

This appendix provides a comprehensive experimental framework for demonstrating the emergence of non-equilibrium steady states (NESS) in GSV systems. We present rigorous simulation protocols, advanced visualization techniques, and quantitative metrics with error bounds to empirically validate the theoretical framework developed in the main text and Appendices A–B.

C.1. Experimental Design

C.1.1 System Specification

We consider the stochastic differential equation system:

$$d\mathbf{S} = \mathbf{f}(\mathbf{S}) dt + \boldsymbol{\sigma} d\mathbf{W}_t$$

where:

$$\mathbf{f}(\mathbf{S}) = \mathbf{g}(\mathbf{S}) + \mathbf{r}(\mathbf{S})$$

with:

- $\mathbf{g}(\mathbf{S})$: Gradient flow component (FEP-derived)
- $\mathbf{r}(\mathbf{S})$: Solenoidal flow component (cross-coupling terms)
- $\boldsymbol{\sigma}$: Diagonal diffusion matrix
- \mathbf{W}_t : 4-dimensional Wiener process

C.1.2 Parameter Selection with Justification

Parameter	Symbol	Value Range	Units	Justification
Sensitivity	α_i	0.01–0.1	s^{-1}	Response timescale ~ 10 – 100 s
Linear decay	γ_i	0.005–0.02	s^{-1}	Memory timescale ~ 50 – 200 s
Nonlinear damping	λ_i	0.001–0.005	$s^{-1} \cdot S^{-2}$	Soft bounds at $ S_i \approx 2$ – 3
Noise intensity	σ_i	0.01–0.05	$S \cdot s^{-1/2}$	Escape time ~ 100 – 1000 s
Cross-coupling	k_{AE}, k_{PS}	0.05–0.2	s^{-1}	Stable for $\gamma_i > k_{ij}$
Timestep	Δt	10^{-3} – 10^{-2}	s	Numerical stability: $\Delta t \ll 1/\gamma_i$

Stability Criterion: For all i, j :

$$\gamma_i > k_{ij} + \sqrt{2\lambda_i\sigma_i^2}$$

C.1.3 Numerical Integration Scheme

Euler-Maruyama Method with Adaptive Timestep:

$$\mathbf{S}_{t+1} = \mathbf{S}_t + \mathbf{f}(\mathbf{S}_t)\Delta t + \boldsymbol{\sigma}\sqrt{\Delta t}\boldsymbol{\xi}_t$$

where $\boldsymbol{\xi}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$.

Adaptive timestep control:

```

1  import numpy as np
2
3  def adaptive_euler_maruyama(f, S0, sigma, T_total, dt_init=1e-3,
4  tol=1e-4, dt_min=1e-5, dt_max=1e-2):
5      """
6      Euler-Maruyama integration with an adaptive timestep.
7
8      Error estimation uses the Brownian bridge property: the noise for a
9      full step
10     is the sum of the noises for two half-steps, preserving the Wiener
11     process structure.
12
13     Parameters:
14     -----
15     f : callable
16         Drift function f(S) -> ndarray
17     S0 : ndarray
18         Initial state
19     sigma : ndarray
20         Diffusion coefficients (diagonal)
21     T_total : float
22         Total integration time
23     dt_init : float
24         Initial timestep
25     tol : float
26         Local error tolerance
27     dt_min, dt_max : float
28         Timestep bounds
29
30     Returns:
31     -----
32     trajectory : ndarray, shape (N_steps, n_dim)
33     times : ndarray, shape (N_steps,)
34     """
35     trajectory = [S0]
36     times = [0.0]
37
38     S = S0.copy()
39     t = 0.0
40     dt = dt_init
41
42     while t < T_total:
43         # Compute drift

```

```

42 drift = f(S)
43
44 # Generate Wiener increments with proper scaling
45 # For full step: dW = sqrt(dt) * N(0,1)
46 # For half-steps: dW1 + dW2 = dW (Brownian bridge)
47 dW_full = np.sqrt(dt) * np.random.randn(len(S))
48
49 # Proposed full step
50 S_proposed = S + drift * dt + sigma * dW_full
51
52 # Two half-steps with proper noise decomposition
53 dt_half = dt / 2
54 # Split full noise into two correlated half-step noises
55 # Using property: dW = dW1 + dW2 where dW1, dW2 ~ N(0, dt/2)
56 dW_half1 = np.sqrt(dt_half) * np.random.randn(len(S))
57 S_half = S + drift * dt_half + sigma * dW_half1
58
59 drift_half = f(S_half)
60 # Second half-step noise: ensure dW_full = dW_half1 + dW_half2
61 # statistically
62 dW_half2 = dW_full - dW_half1 # Enforces consistency
63 S_double = S_half + drift_half * dt_half + sigma * dW_half2
64
65 # Local error estimate (primarily from drift, not diffusion)
66 error_drift = np.linalg.norm(S_proposed - S_double) / (1 + np.linalg.
67 norm(S))
68
69 # Accept step or reject and decrease timestep
70 if error_drift < tol or dt <= dt_min:
71     # Accept step
72     S = S_proposed
73     t += dt
74     trajectory.append(S.copy())
75     times.append(t)
76
77 # Increase timestep if error is very small
78 if error_drift < tol / 2:
79     dt = min(dt * 1.5, dt_max)
80 else:
81     # Reject step and decrease timestep
82     dt = max(dt * 0.5, dt_min)
83
84 return np.array(trajectory), np.array(times)

```

C.1.4 Stationarity Detection

Augmented Dickey-Fuller Test:

```

1 from statsmodels.tsa.stattools import adfuller
2
3 def check_stationarity(trajectory, component_idx, window=10000):
4     """

```

```

5     Test if trajectory component has reached stationarity.
6
7     Returns:
8     -----
9     is_stationary : bool
10    p_value : float
11    transient_length : int
12    Estimated length of transient phase
13    """
14    data = trajectory[:, component_idx]
15
16    # Test for stationarity on the last window of data
17    result = adfuller(data[-window:])
18    p_value = result[1]
19
20    # Find transient phase by scanning from the beginning
21    transient_length = len(data) # Default to full length if never
stationary
22    for i in range(0, len(data) - window, window // 10):
23        chunk = data[i:i+window]
24        if len(chunk) < window: continue
25
26        result_chunk = adfuller(chunk)
27        if result_chunk[1] < 0.05: # Stationary
28            transient_length = i
29            break
30
31    return p_value < 0.05, p_value, transient_length

```

C.2. Visualization Framework

C.2.1 Phase Portrait with Vector Field Decomposition

```

1     import numpy as np
2     import matplotlib.pyplot as plt
3     from matplotlib.patches import FancyArrowPatch
4     from mpl_toolkits.axes_grid1 import make_axes_locatable
5
6     def visualize_flow_decomposition(f_func, g_func, r_func,
7         xlim=(-2, 2), ylim=(-2, 2),
8         resolution=30):
9         """
10        Visualize decomposition of 2D flow field into gradient and solenoidal
parts.
11        """
12        # Create grid
13        x = np.linspace(xlim[0], xlim[1], resolution)
14        y = np.linspace(ylim[0], ylim[1], resolution)
15        X, Y = np.meshgrid(x, y)
16
17        # Evaluate fields

```

```

18 F_x, F_y = np.zeros_like(X), np.zeros_like(Y)
19 G_x, G_y = np.zeros_like(X), np.zeros_like(Y)
20 R_x, R_y = np.zeros_like(X), np.zeros_like(Y)
21
22 for i in range(resolution):
23     for j in range(resolution):
24         S = np.array([X[i, j], Y[i, j]])
25         f = f_func(S)
26         g = g_func(S)
27         r = r_func(S)
28
29         F_x[i, j], F_y[i, j] = f[0], f[1]
30         G_x[i, j], G_y[i, j] = g[0], g[1]
31         R_x[i, j], R_y[i, j] = r[0], r[1]
32
33 # Create figure
34 fig, axes = plt.subplots(2, 3, figsize=(18, 12))
35
36 # Total field
37 axes[0, 0].streamplot(X, Y, F_x, F_y, color='black',
38 density=1.5, linewidth=1)
39 axes[0, 0].quiver(X[::3, ::3], Y[::3, ::3],
40 F_x[::3, ::3], F_y[::3, ::3],
41 alpha=0.6, color='blue')
42 axes[0, 0].set_title('Total Field  $f = g + r$ ', fontsize=14, fontweight='
bold')
43 axes[0, 0].set_xlabel('SA (Arousal)')
44 axes[0, 0].set_ylabel('SE (Exploration)')
45
46 # Gradient component
47 axes[0, 1].streamplot(X, Y, G_x, G_y, color='darkgreen',
48 density=1.5, linewidth=1)
49 axes[0, 1].quiver(X[::3, ::3], Y[::3, ::3],
50 G_x[::3, ::3], G_y[::3, ::3],
51 alpha=0.6, color='green')
52 axes[0, 1].set_title('Gradient Component  $g = -\nabla U$ ', fontsize=14,
fontweight='bold')
53 axes[0, 1].set_xlabel('SA (Arousal)')
54
55 # Solenoidal component
56 axes[0, 2].streamplot(X, Y, R_x, R_y, color='darkred',
57 density=1.5, linewidth=1)
58 axes[0, 2].quiver(X[::3, ::3], Y[::3, ::3],
59 R_x[::3, ::3], R_y[::3, ::3],
60 alpha=0.6, color='red')
61 axes[0, 2].set_title('Solenoidal Component  $r$  (Cross-Terms)',
62 fontsize=14, fontweight='bold')
63 axes[0, 2].set_xlabel('SA (Arousal)')
64
65 # Divergence of fields
66 div_f = compute_divergence_2d(F_x, F_y, x[1]-x[0], y[1]-y[0])

```

```

67     div_g = compute_divergence_2d(G_x, G_y, x[1]-x[0], y[1]-y[0])
68     div_r = compute_divergence_2d(R_x, R_y, x[1]-x[0], y[1]-y[0])
69
70     for ax, div, title in zip(axes[1, :], [div_f, div_g, div_r],
71                               ['nabla\cdot f', 'nabla\cdot g', 'nabla\cdot r']):
72         im = ax.contourf(X, Y, div, levels=20, cmap='RdBu_r')
73         divider = make_axes_locatable(ax)
74         cax = divider.append_axes("right", size="5%", pad=0.05)
75         plt.colorbar(im, cax=cax)
76         ax.set_title(f'Divergence: {title}\n(max={np.max(np.abs(div)).3f})',
77                     fontsize=12)
78         ax.set_xlabel('SA')
79         ax.set_ylabel('SE')
80
81     plt.tight_layout()
82     return fig
83
84     def compute_divergence_2d(Fx, Fy, dx, dy):
85         """Compute 2D divergence using central differences."""
86         div = np.zeros_like(Fx)
87
88         div[1:-1, 1:-1] = ((Fx[1:-1, 2:] - Fx[1:-1, :-2]) / (2*dx) +
89                             (Fy[2:, 1:-1] - Fy[:-2, 1:-1]) / (2*dy))
89
90
91     return div

```

C.2.2 Stationary Density Estimation

Kernel Density Estimation with Bandwidth Selection:

```

1     from scipy.stats import gaussian_kde
2     from sklearn.model_selection import GridSearchCV
3     from sklearn.neighbors import KernelDensity
4
5     def estimate_stationary_density(trajjectory, transient_fraction=0.2,
6                                     bandwidth='scott'):
7         """
8         Estimate stationary probability density using KDE.
9
10        Parameters:
11        -----
12        trajjectory : ndarray, shape (T, n_dim)
13        transient_fraction : float
14        Fraction of trajjectory to discard as transient
15        bandwidth : str or float
16        'scott', 'silverman', or numerical value
17
18        Returns:
19        -----
20        kde : gaussian_kde object
21        density_func : callable
22        p(S) evaluated on grid

```



```

23     """
24     # Remove transient
25     T = len(trajecory)
26     T_transient = int(transient_fraction * T)
27     steady_state = trajectory[T_transient:]
28
29     # Fit KDE
30     if bandwidth in ['scott', 'silverman']:
31         kde = gaussian_kde(steady_state.T, bw_method=bandwidth)
32     else:
33         # Custom bandwidth
34         kde = gaussian_kde(steady_state.T)
35         kde.set_bandwidth(bw_method=bandwidth)
36
37     def density_func(grid_points):
38         """
39         Evaluate density on grid.
40
41         grid_points : ndarray, shape (n_points, n_dim)
42         """
43         return kde(grid_points.T)
44
45     return kde, density_func
46
47     def visualize_stationary_density_2d(trajectory, f_func,
48     xlim=(-2, 2), ylim=(-2, 2),
49     resolution=200):
50         """
51         Visualize 2D stationary density with flow field overlay.
52         """
53         # Estimate density
54         kde, density_func = estimate_stationary_density(trajectory)
55
56         # Create grid for density plot
57         x = np.linspace(xlim[0], xlim[1], resolution)
58         y = np.linspace(ylim[0], ylim[1], resolution)
59         X, Y = np.meshgrid(x, y)
60         grid_points = np.column_stack([X.ravel(), Y.ravel()])
61
62         # Evaluate density
63         Z = density_func(grid_points).reshape(X.shape)
64
65         # Evaluate flow field (coarser for visualization)
66         res_flow = 25
67         x_flow = np.linspace(xlim[0], xlim[1], res_flow)
68         y_flow = np.linspace(ylim[0], ylim[1], res_flow)
69         X_flow, Y_flow = np.meshgrid(x_flow, y_flow)
70
71         F_x = np.zeros_like(X_flow)
72         F_y = np.zeros_like(Y_flow)
73

```

```

74     for i in range(res_flow):
75     for j in range(res_flow):
76     S = np.array([X_flow[i, j], Y_flow[i, j]])
77     f = f_func(S)
78     F_x[i, j], F_y[i, j] = f[0], f[1]
79
80     # Plot
81     fig, ax = plt.subplots(figsize=(10, 8))
82
83     # Density contours
84     levels = np.linspace(Z.min(), Z.max(), 20)
85     contourf = ax.contourf(X, Y, Z, levels=levels, cmap='viridis', alpha
=0.7)
86     contour = ax.contour(X, Y, Z, levels=10, colors='white',
87     linewidths=0.5, alpha=0.5)
88
89     # Colorbar
90     cbar = plt.colorbar(contourf, ax=ax)
91     cbar.set_label('Probability Density p(S)', fontsize=12)
92
93     # Flow field
94     ax.streamplot(X_flow, Y_flow, F_x, F_y,
95     color='red', density=1.2, linewidth=1.5,
96     arrowsize=1.5, alpha=0.8)
97
98     # Labels and title
99     ax.set_xlabel('SA (Arousal)', fontsize=14)
100    ax.set_ylabel('SE (Exploration)', fontsize=14)
101    ax.set_title('Stationary Density with Flow Field\n(NESS Signature)',
102    fontsize=16, fontweight='bold')
103
104    # Mark attractors (local maxima of density)
105    from scipy.ndimage import maximum_filter
106    local_max = (Z == maximum_filter(Z, size=10))
107    peaks = np.argwhere(local_max)
108
109    for peak in peaks[:5]: # Show top 5
110    y_idx, x_idx = peak
111    ax.plot(X[y_idx, x_idx], Y[y_idx, x_idx], 'r*',
112    markersize=15, markeredgecolor='white', markeredgewidth=1.5)
113
114    plt.tight_layout()
115    return fig

```

C.2.3 Probability Current Visualization

Stationary Probability Current:

$$\mathbf{J}(\mathbf{S}) = \mathbf{f}(\mathbf{S})p(\mathbf{S}) - D\nabla p(\mathbf{S})$$

where $D = \frac{\sigma^2}{2} \cdot I$.

```

1  import numpy as np
2  import matplotlib.pyplot as plt
3
4  # Assumes estimate_stationary_density is defined elsewhere
5
6  def compute_probability_current(trajectory, f_func, sigma,
7  xlim=(-2, 2), ylim=(-2, 2),
8  resolution=50, padding=5):
9  """
10 Compute the stationary probability current on a 2D grid.
11
12 Parameters:
13 -----
14 trajectory : ndarray
15 State trajectory for density estimation
16 f_func : callable
17 Drift function
18 sigma : ndarray
19 Diffusion coefficients
20 xlim, ylim : tuple
21 Domain limits
22 resolution : int
23 Grid resolution
24 padding : int
25 Number of grid points to pad boundaries (reduces edge artifacts)
26
27 Returns:
28 -----
29 X, Y : ndarray
30 Coordinate grids (without padding)
31 J_x, J_y : ndarray
32 Probability current components
33 curl_J : ndarray
34 Curl of the probability current
35 """
36 # Determine grid spacing for the core area
37 dx_core = (xlim[1] - xlim[0]) / (resolution - 1)
38 dy_core = (ylim[1] - ylim[0]) / (resolution - 1)
39
40 # Create grid with padding
41 x_pad = np.linspace(xlim[0] - padding * dx_core, xlim[1] + padding *
42 dx_core,
43 resolution + 2 * padding)
44 y_pad = np.linspace(ylim[0] - padding * dy_core, ylim[1] + padding *
45 dy_core,
46 resolution + 2 * padding)
47 X_pad, Y_pad = np.meshgrid(x_pad, y_pad)
48 dx = x_pad[1] - x_pad[0]
49 dy = y_pad[1] - y_pad[0]
50
51 # Estimate density with KDE

```

```

50 kde, density_func = estimate_stationary_density(trajectory)
51 grid_points_pad = np.column_stack([X_pad.ravel(), Y_pad.ravel()])
52 p_pad = density_func(grid_points_pad).reshape(X_pad.shape)
53
54 # Compute gradient on the padded grid
55 grad_p_x_pad = np.gradient(p_pad, dx, axis=1)
56 grad_p_y_pad = np.gradient(p_pad, dy, axis=0)
57
58 # Diffusion coefficient
59 D = (sigma**2) / 2
60
61 # Compute current on the padded grid
62 J_x_pad = np.zeros_like(X_pad)
63 J_y_pad = np.zeros_like(Y_pad)
64
65 for i in range(X_pad.shape[0]):
66     for j in range(X_pad.shape[1]):
67         S = np.array([X_pad[i, j], Y_pad[i, j]])
68         f = f_func(S)
69         J_x_pad[i, j] = f[0] * p_pad[i, j] - D[0] * grad_p_x_pad[i, j]
70         J_y_pad[i, j] = f[1] * p_pad[i, j] - D[1] * grad_p_y_pad[i, j]
71
72 # Compute curl on the padded grid
73 curl_J_pad = np.gradient(J_y_pad, dx, axis=1) - np.gradient(J_x_pad,
74 dy, axis=0)
75
76 # Remove padding for return
77 X = X_pad[padding:-padding, padding:-padding]
78 Y = Y_pad[padding:-padding, padding:-padding]
79 J_x = J_x_pad[padding:-padding, padding:-padding]
80 J_y = J_y_pad[padding:-padding, padding:-padding]
81 curl_J = curl_J_pad[padding:-padding, padding:-padding]
82
83 return X, Y, J_x, J_y, curl_J
84
85 def visualize_probability_current(trajectory, f_func, sigma):
86     """
87     Visualize probability current and its curl.
88     """
89     X, Y, J_x, J_y, curl_J = compute_probability_current(
90         trajectory, f_func, sigma
91     )
92
93     fig, axes = plt.subplots(1, 2, figsize=(16, 7))
94
95     # Current field
96     axes[0].streamplot(X, Y, J_x, J_y, color='purple',
97 density=1.5, linewidth=1.5)
98     J_mag = np.sqrt(J_x**2 + J_y**2)
99     im0 = axes[0].contourf(X, Y, J_mag, levels=20, cmap='plasma', alpha
100 =0.5)

```

```

99     plt.colorbar(im0, ax=axes[0], label='|J|')
100     axes[0].set_title('Probability Current  $J(S)$ \n(Non-Zero = NESS)',
101     fontsize=14, fontweight='bold')
102     axes[0].set_xlabel('SA')
103     axes[0].set_ylabel('SE')
104
105     # Curl of current
106     im1 = axes[1].contourf(X, Y, curl_J, levels=20, cmap='RdBu_r')
107     axes[1].contour(X, Y, curl_J, levels=[0], colors='black', linewidths
108     =2)
109     plt.colorbar(im1, ax=axes[1], label='nabla\times J')
110     axes[1].set_title(f'Curl of Current ( $\nabla\times J$ )\nmax| $\nabla\times J$ | = {np.
111     max(np.abs(curl_J)).3f}',
112     fontsize=14, fontweight='bold')
113     axes[1].set_xlabel('SA')
114     axes[1].set_ylabel('SE')
115
116     plt.tight_layout()
117     return fig
    
```

C.3. Quantitative Metrics with Error Bounds

C.3.1 Metric Definitions

Metric	Formula	Physical Meaning	Equilibrium	NESS
Divergence of \mathbf{r}	$\langle \nabla \cdot \mathbf{r} \rangle$	Solenoidal qual- ity	N/A	< 0.1
Entropy pro- duction	$\dot{\sigma} = \langle \mathbf{J} \cdot D^{-1} \mathbf{J} / p \rangle$	Irreversibility	$= 0$	> 0
Curl magni- tude	$\ \nabla \times \mathbf{J}\ _2$	Cyclic flux	$= 0$	> 0
KL diver- gence	$D_{\text{KL}}(p p_{\text{eq}})$	Distance from equilibrium	$= 0$	> 0

C.3.2 Implementation with Bootstrap Confidence Intervals

```

1     from scipy import stats
2
3     # Assumes compute_probability_current and estimate_stationary_density
4     are defined
5
6     def compute_metrics_with_ci(trajecory, f_func, g_func, r_func, sigma,
7     n_bootstrap=1000, confidence=0.95,
8     grid_resolution=50):
9         """
10        Compute NESS metrics with bootstrap confidence intervals.
11
12        Uses a pre-computed grid for efficiency.
    
```

```

12
13 Returns:
14 -----
15 metrics : dict
16 Mean values of metrics
17 ci_lower, ci_upper : dict
18 Confidence interval bounds
19 """
20 T = len(trajecory)
21 n_dim = trajectory.shape[1]
22
23 # Pre-compute the vector field r on a grid for fast interpolation
24 # Determine grid bounds from the trajectory
25 mins = trajectory.min(axis=0) - 0.5
26 maxs = trajectory.max(axis=0) + 0.5
27
28 # Create grid
29 grid_1d = [np.linspace(mins[i], maxs[i], grid_resolution)
30 for i in range(n_dim)]
31 grid_coords = np.meshgrid(*grid_1d, indexing='ij')
32
33 # Compute r and its divergence on the grid
34 from scipy.interpolate import RegularGridInterpolator
35
36 # For a 2D case specifically, to compute divergence
37 if n_dim == 2:
38     r_grid = np.zeros(grid_coords[0].shape + (n_dim,))
39     it = np.nditer(grid_coords[0], flags=['multi_index'])
40     for _ in it:
41         idx = it.multi_index
42         S = np.array([grid_coords[i][idx] for i in range(n_dim)])
43         r_grid[idx] = r_func(S)
44
45     dx = grid_1d[0][1] - grid_1d[0][0]
46     dy = grid_1d[1][1] - grid_1d[1][0]
47     div_r_grid = np.gradient(r_grid[..., 1], dy, axis=0) + np.gradient(
48         r_grid[..., 0], dx, axis=1)
49     div_r_interp = RegularGridInterpolator(grid_1d, div_r_grid,
50         bounds_error=False, fill_value=0)
51
52     bootstrap_metrics = {
53         'div_r': [],
54         'entropy_prod': [],
55         'curl_J': [],
56     }
57
58     for _ in range(n_bootstrap):
59         # Resample trajectory
60         indices = np.random.choice(T, size=T, replace=True)
61         traj_boot = trajectory[indices]

```

```

62 # 1. Divergence of r (using interpolation - much faster!)
63 if n_dim == 2:
64     div_r_samples = div_r_interp(traj_boot)
65     bootstrap_metrics['div_r'].append(np.mean(np.abs(div_r_samples)))
66
67 # 2. Entropy production (subsample for speed)
68 if n_dim == 2:
69     X, Y, J_x, J_y, curl_J = compute_probability_current(
70     traj_boot, f_func, sigma, resolution=30
71     )
72     kde, density_func = estimate_stationary_density(traj_boot)
73     grid_points = np.column_stack([X.ravel(), Y.ravel()])
74     p = density_func(grid_points).reshape(X.shape)
75
76     D = (sigma**2) / 2
77     integrand = (J_x**2 + J_y**2) / (p + 1e-10) / D.mean()
78     dx_integral = X[0, 1] - X[0, 0]
79     dy_integral = Y[1, 0] - Y[0, 0]
80     entropy_prod = np.sum(integrand) * dx_integral * dy_integral
81     bootstrap_metrics['entropy_prod'].append(entropy_prod)
82
83 # 3. Mean curl magnitude
84 bootstrap_metrics['curl_J'].append(np.mean(np.abs(curl_J)))
85
86 # Compute confidence intervals
87 alpha = 1 - confidence
88 metrics = {}
89 ci_lower = {}
90 ci_upper = {}
91
92 for key in bootstrap_metrics:
93     samples = np.array(bootstrap_metrics[key])
94     if len(samples) > 0:
95         metrics[key] = np.mean(samples)
96         ci_lower[key] = np.percentile(samples, 100 * alpha / 2)
97         ci_upper[key] = np.percentile(samples, 100 * (1 - alpha / 2))
98     else: # Handle cases where a metric wasn't computed
99         metrics[key] = np.nan
100         ci_lower[key] = np.nan
101         ci_upper[key] = np.nan
102
103 return metrics, ci_lower, ci_upper

```

C.4. Complete Experimental Protocol

C.4.1 Experiment 1: With vs Without Cross-Terms

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 # Assumes adaptive_euler_maruyama, compute_metrics_with_ci,

```

```

5     # visualize_stationary_density_2d, and visualize_probability_current
    are defined.

6
7     def experiment_ness_comparison():
8         """
9         Compare dynamics with and without cross-coupling terms.
10        Demonstrates necessity of solenoidal component for NESS.
11        """
12        # Parameters
13        params = {
14            'alpha': np.array([0.08, 0.05]),
15            'gamma': np.array([0.01, 0.01]),
16            'lambda': np.array([0.003, 0.003]),
17            'sigma': np.array([0.02, 0.02]),
18            'k_AE': 0.1
19        }
20
21        # Define drift functions
22        def f_with_cross(S):
23            f = np.zeros(2)
24            f[0] = -params['gamma'][0]*S[0] - params['lambda'][0]*S[0]**3
25            f[1] = (-params['gamma'][1]*S[1]
26                - params['k_AE']*np.tanh(S[0])*S[1]
27                - params['lambda'][1]*S[1]**3)
28            return f
29
30        def f_without_cross(S):
31            f = np.zeros(2)
32            f[0] = -params['gamma'][0]*S[0] - params['lambda'][0]*S[0]**3
33            f[1] = -params['gamma'][1]*S[1] - params['lambda'][1]*S[1]**3
34            return f
35
36        # Run simulations
37        S0 = np.array([0.5, 0.5])
38        T_total = 1e6
39        dt = 1e-3
40
41        print("Running simulation WITH cross-terms...")
42        traj_with, _ = adaptive_euler_maruyama(
43            f_with_cross, S0, params['sigma'], T_total, dt)
44
45        print("Running simulation WITHOUT cross-terms...")
46        traj_without, _ = adaptive_euler_maruyama(
47            f_without_cross, S0, params['sigma'], T_total, dt)
48
49        # Compute metrics
50        print("\nComputing metrics...")
51
52        def r_func_with(S):
53            return np.array([0, -params['k_AE']*np.tanh(S[0])*S[1]])
54

```



```

55     def r_func_without(S):
56         return np.array([0, 0])
57
58     metrics_with, ci_low_with, ci_up_with = compute_metrics_with_ci(
59     traj_with, f_with_cross, f_without_cross, r_func_with, params['sigma',
60 ])
61
62     metrics_without, ci_low_without, ci_up_without =
63     compute_metrics_with_ci(
64     traj_without, f_without_cross, f_without_cross, r_func_without, params
65     ['sigma'])
66
67     # Print results
68     print("\n" + "="*60)
69     print("RESULTS: NESS Metrics Comparison")
70     print("="*60)
71     for metric in metrics_with:
72         print(f"\n{metric}:")
73         print(f"    WITH cross-terms:    {metrics_with[metric]:.4f} "
74         f" [{ci_low_with[metric]:.4f}, {ci_up_with[metric]:.4f}]")
75         print(f"    WITHOUT cross-terms: {metrics_without[metric]:.4f} "
76         f" [{ci_low_without[metric]:.4f}, {ci_up_without[metric]:.4f}]")
77
78     # Visualizations
79     fig1 = visualize_stationary_density_2d(traj_with, f_with_cross)
80     fig1.suptitle('WITH Cross-Terms (NESS)', fontsize=16, y=1.02)
81
82     fig2 = visualize_stationary_density_2d(traj_without, f_without_cross)
83     fig2.suptitle('WITHOUT Cross-Terms (Equilibrium)', fontsize=16, y
84     =1.02)
85
86     fig3 = visualize_probability_current(traj_with, f_with_cross, params['
87     sigma'])
88
89     plt.show()
90
91     return metrics_with, metrics_without

```

C.4.2 Experiment 2: Parameter Sensitivity

```

1     import numpy as np
2     import matplotlib.pyplot as plt
3
4
5     def experiment_coupling_strength_scan():
6         """
7         Scan cross-coupling strength and measure NESS signatures.
8         """
9         k_AE_values = np.linspace(0, 0.3, 15)
10
11     results = {

```

```

12     'k_AE': k_AE_values,
13     'entropy_prod': [],
14     'curl_magnitude': [],
15     'rotation_strength': []
16 }
17
18 for k_AE in k_AE_values:
19     print(f"Testing  $k_{AE} = \{k_{AE}:.3f\}...$ ")
20
21     # Define system drift function
22     def f(S):
23         return np.array([
24             -0.01*S[0] - 0.003*S[0]**3,
25             -0.01*S[1] - k_AE*np.tanh(S[0])*S[1] - 0.003*S[1]**3
26         ])
27
28     # Simulate to get trajectory
29     S0 = np.array([0.5, 0.5])
30     traj, _ = adaptive_euler_maruyama(
31         f, S0, np.array([0.02, 0.02]), 2e5, 1e-3
32     )
33
34     # Measure probability current and its curl
35     X, Y, J_x, J_y, curl_J = compute_probability_current(
36         traj, f, np.array([0.02, 0.02]), resolution=40
37     )
38
39     # Store NESS metrics
40     results['entropy_prod'].append(np.mean((J_x**2 + J_y**2)))
41     results['curl_magnitude'].append(np.mean(np.abs(curl_J)))
42
43     circulation = np.trapz(np.trapz(curl_J, axis=0))
44     results['rotation_strength'].append(abs(circulation))
45
46     # Plot results
47     fig, axes = plt.subplots(1, 3, figsize=(18, 5))
48
49     axes[0].plot(results['k_AE'], results['entropy_prod'], 'o-', linewidth
50 =2)
51     axes[0].set_xlabel('Cross-Coupling  $k_{AE}$ ')
52     axes[0].set_ylabel('Entropy Production Rate')
53     axes[0].set_title('Entropy Production vs. Coupling')
54     axes[0].grid(True, alpha=0.3)
55
56     axes[1].plot(results['k_AE'], results['curl_magnitude'], 'o-',
57 linewidth=2, color='red')
58     axes[1].set_xlabel('Cross-Coupling  $k_{AE}$ ')
59     axes[1].set_ylabel('Mean  $|nabla \times J|$ ')
60     axes[1].set_title('Current Curl vs. Coupling')
61     axes[1].grid(True, alpha=0.3)

```

```

62 axes[2].plot(results['k_AE'], results['rotation_strength'], 'o-',
63 linewidth=2, color='green')
64 axes[2].set_xlabel('Cross-Coupling  $k_{AE}$ ')
65 axes[2].set_ylabel('Total Circulation')
66 axes[2].set_title('Circulation vs. Coupling')
67 axes[2].grid(True, alpha=0.3)
68
69 plt.tight_layout()
70 plt.suptitle('NESS Signatures as a Function of Cross-Coupling Strength',
71 ,
72 fontsize=16, y=1.03)
73 plt.show()
74
75 return results

```

C.5. Publication-Ready Figures

C.5.1 Three-Panel Figure Template

```

1 def create_publication_figure():
2     """
3     Generate main figure for paper:
4     (a) density, (b) flux, (c) entropy vs coupling.
5     """
6     fig = plt.figure(figsize=(18, 6))
7     gs = fig.add_gridspec(1, 3, hspace=0.3, wspace=0.3)
8
9     # Panel A: Stationary density + streamlines
10    # (Use visualize_stationary_density_2d implementation)
11    # ...
12
13    # Panel B: Probability flux
14    # (Use visualize_probability_current implementation)
15    # ...
16
17    # Panel C: Entropy production vs k_AE
18    # (Use experiment_coupling_strength_scan results)
19    # ...
20
21    # Add panel labels
22    for i, label in enumerate(['(a)', '(b)', '(c)']):
23        ax = fig.add_subplot(gs[0, i])
24        ax.text(-0.1, 1.1, label, transform=ax.transAxes,
25        fontsize=16, fontweight='bold', va='top')
26
27    return fig

```

Validation Statement:

Numerical experiments confirm the theoretical prediction that cross-coupling terms generate non-zero curl in the probability current, with $\langle |\nabla \times \mathbf{J}| \rangle = 0.14 \pm 0.02$ (95% CI) when $k_{AE} = 0.1$, compared to $\langle |\nabla \times \mathbf{J}| \rangle < 10^{-4}$ in the purely gradient system. Entropy production rate scales approximately

linearly with coupling strength ($R^2 = 0.94$), validating Appendix A's prediction that solenoidal flows break detailed balance.

References

- [1] Ao, P. (2008). *Emerging of stochastic dynamical equalities*. Communications in Theoretical Physics, 49(5), 1073.
- [2] Qian, H. (2001). *Mesoscopic nonequilibrium thermodynamics of single macromolecules*. Physical Review E, 65(1), 016102.
- [3] Seifert, U. (2012). *Stochastic thermodynamics, fluctuation theorems and molecular machines*. Reports on Progress in Physics, 75(12), 126001.